

DeepSeek Harness

白皮书

From zero to one with DeepSeek Harness

Version 0.1.0-rc.6 · 2026-08-13 · 15 chapters + Appendices

目录

Chapter 1: Understanding DeepSeek Harness

> *Goal: build a complete mental model of DeepSeek Harness (dsh) from absolute zero — **what it is, why it matters, how it differs from mainstream agents, and when to use it**. No prior knowledge required.*

1.1 Three intuitions first

① What is dsh? — "the LEGO baseplate for agents."

LEGO gives you the baseplate and standard bricks (runtime + core plugins); you assemble freely (add plugins, swap UIs, change behavior). Claude Code, by contrast, is "a finished car" — great to drive, but modifying the engine means asking the vendor.

② Why "harness"? — the engineering layer around a model.

A model (DeepSeek V4) only "replies with text." To make it work in your repository (read files, run commands, edit code, loop), you need an engineering layer on top: session management, tool calling, context control, error recovery. **That layer is a harness.** dsh is DeepSeek open-sourcing that layer.

③ Why open-source it now? — agents entered the "programmable era".

2025 was model-capability competition; 2026 is agent-engineering competition. Open-sourcing the harness is the strategic move to make "how to organize agents" an open ecosystem — like Android did for phones.

1.2 Official facts


One-liner: DeepSeek Harness (dsh) is DeepSeek's open-source agent runtime with an "everything is a plugin" architecture, built on the Cordis plugin container.

Fact	Value	
Open-sourced	2026-08-13	
License	MIT	
Language	TypeScript Node.js ≥ 22	
Version	0.1.0-rc.x rc.6; fast iteration, breaking changes expected	
Runtime	Cordis	
Built-in profiles	web + headless	

1.3 How "everything is a plugin" works

```
flowchart TB
    subgraph Profile["Your profile (bootable form)"]
        P1["dsh web (Web UI)"]
        P2["dsh headless (CLI)"]
        P3["Custom profile (TUI/desktop/bot...)"]
    end
    subgraph Plugins["Capability layer (each capability = one plugin)"]
        L["llm: model access + reasoning effort"]
        T["tools: files/terminal/search/skills"]
        S["session: conversation persistence"]
        C["client: UI (web/terminal)"]
        ST["settings: user config"]
    end
    subgraph Cordis["Cordis plugin container"]
        DI["DI · events · lifecycle"]
    end
    Profile --> Cordis
    Cordis --> Plugins
    Plugins --> L
    Plugins --> T
    Plugins --> S
    Plugins --> C
    Plugins --> ST
```

Layered view



Three concepts you must know

- ① Profile (bootable form)** — a directory `$DSH_HOME/profiles/<name>/` with `package.json` (plugin deps + manifest) and `cordis.patch.yml` (your patch layer). Load order: built-in bundles → profile patch → global patch → `--patch` overlays.
- ② host half / client half** — one npm package can carry both: the host half (`apply(ctx)` on Node: tools, services, events) and the client half (browser UI, declared via `dsh.client` in `package.json`).
- ③ Extension points** — official principle: *"Plugins, not loop changes."* Use hooks, don't fork the core. Common ones: `agent/request` waterfall (change request config per step), `conversationEvents.register`, `ctx.slots.inject` (inject UI), `settings` service.

1.4 dsh vs mainstream agents

Capability matrix

Dimension	dsh	Claude Code	OpenAI Codex	OpenCode	Gemini CLI	Kimi CLI	
Open source	✔ MIT	✗	✗	✔ MIT	✗	✗	
Model binding	model-agnostic	Claude	GPT	any	Gemini	Kimi	
Official runtime	✔ + plugin ecosystem	product	product	client only	product	product	

Plugin system	first-class, 60+ official pkgs	config/hooks	config	config	none	none	
Custom UI	✔ client half	✗	✗	partial	✗	✗	
Automation/CI	✔ headless	✔	✔	✔	✔	✔	
TUI	plugin-provided	✔ built-in	✔ built-in	✔ built-in	✔	✔	
Ecosystem stage	day-zero 2026-08-13	mature	mature	mature	mature	early	
Best for	deep customization + ecosystem	out-of-box	out-of-box	OpenCode users	Google	Kimi	

Case: same task, different doors

****Task**:** "Find every call of function X in the repo and change one argument."

Agent	How you do it	
dsh web	dsh web → type → model uses Grep/Read/Edit tools; add plugin sidebars Git panel, speed-up plugin	
dsh headless	dsh --profile headless "task" → prints result, exits — CI-friendly	
Claude Code / Codex / OpenCode / Kimi	open the TUI → type → model does it	

****The difference**:** same sentence, but dsh gives you a choice dimension — swap the UI and toolchain. Others ship a fixed interface.

Case: real workloads (our measurements, 2026-08-13)

Scenario	dsh measured	Note	
Simple file create	cold start ~110s first round, context injection → hot cache ~1s	reasoning effort is the main variable	
50-step tool chain	LLM 10m+, tool calls	per-step thinking	

	9m+	accumulates — that's the speed plugin's value	
Plugin dev	zero to runnable plugin: 1 day tests + live verify	clear extension points	
vs Claude Code same task	dsh + V4-Flash $\approx$ 1/10–1/30 of Claude cost	price dimension favors dsh ecosystem	

1.5 When to use dsh

****✅ Adopt now****: model-agnostic + UI-optional + behavior-mutable base; be an early ecosystem contributor; run agents on servers/CI; cost-sensitive.

****🛑 Wait****: you need an out-of-the-box coding assistant; can't accept rc breaking changes; deeply dependent on a vendor's exclusive model feature.

1.6 FAQ

- ****Is dsh a model?**** No. It's a runtime; models plug in via the `llm` plugin (DeepSeek V4 native, OpenAI-compatible others possible).
- ****dsh vs OpenCode**** Both open-source agent clients; OpenCode is "client + config", dsh is "runtime + official plugin ecosystem" (official bundles, 60+ packages).
- ****Do I need TypeScript**** To *use* it, no. To *write plugins* (Ch.4), basic TS — with full code in the handbook.
- ****Is it stable**** rc stage; use for ecosystem/tinkering now, production core after 0.1.0.
- ****Why learn dsh now**** Day-zero ecosystem + official encouragement + Chinese-tutorial vacuum — first-mover window.

****Next****: [Chapter 2: 5-Minute Quickstart](./02-quickstart.md)

Chapter 2: 5-Minute Quickstart

> *Goal: ****follow along and get it running.**** Every command shows its expected output and common errors. Open a terminal and go.*

2.1 Prerequisites (30-second check)

Need	Check	Pass	
Node.js $\geq$ 22	node --version	v22.x+ 24 recommended	
npm	npm --version	prints a version	
Network	npm registry reachable	can install packages	
optional DeepSeek API key	https://platform.deepseek.com	for real conversations	

> *dsh starts without a key (UI opens), but conversations need one. Examples assume it's configured.*

2.2 Install (two ways)

****One-liner (recommended for new users)****

```
npx -y @deepseek-ai/dsh --version
```

First run downloads dsh (large package, 40+ plugin modules, 1-3 min). You're good when you see:

```
<!-- [style] 输出/目录类代码块统一补 text 语言标签 -->
0.1.0-rc.6
```

****Global install (recommended for frequent use)****

```
npm install -g @deepseek-ai/dsh dsh --version
```

2.3 Mode 1: Web UI (`dsh web`)

Start

```
dsh web
```

Expected output:

```
dsh web: http://127.0.0.1:3080
```

Open <http://127.0.0.1:3080> in your browser.

UI map (see screenshot)

Area	What's there	
Left	session list / workspace switch / new session	
Center	chat: input box, model selector DeepSeek V4 Flash, reasoning level High	
Right/bottom	plugin sidebar empty by default; appears after installing community plugins	
Top-right	Session log / Trajectory	

First conversation

- 1. Click ****New session****
- 2. Type: `Hello, introduce yourself in one sentence`
- 3. Press Enter

Expected reply (roughly):

> *Hello! I'm a DeepSeek-powered AI coding assistant...*

Models & reasoning effort

Click the model selector next to the input:

Model	Positioning	
deepseek-v4-flash default	value: fast, cheap, enough for daily work	
deepseek-v4-pro	flagship: stronger, more expensive/slower	

****Reasoning effort**** (three steps since 2026-08-13):

Level	Speed	Quality	Suggested use	
low	fastest	adequate	simple/ deterministic rounds, batch, cheap chain steps	
high default	medium	good	daily agent tasks	
max	slowest	best	hard reasoning, long planning	

> 💡 **Key performance insight**: the model re-thinks **before every tool call**.
 Measured: a "create file" task spends ~90% of wall-clock in thinking; a 50-step tool chain accumulates minutes. **Lowering the reasoning effort is the highest-leverage speedup** (Ch.6 + the example speed-up plugin).

2.4 Mode 2: Headless (one-shot tasks, scripts/CI)

```
dsh --profile headless "Hello, introduce yourself in one sentence"
```

Expected output (prints result, process exits):

```
Hello! I'm a DeepSeek-powered AI coding assistant...
```

Headless value:

- **Automation**: CI, servers, cron
- **Script-friendly**: non-zero exit = failure; output pipes
- **Isolation**: fresh session per call (`--resume` restores; see `dsh --profile headless --help``)

Example: daily digest script:

```
dsh --profile headless "Read today's git log in the workspace and write a Chinese daily summary" > daily-report.md echo "exit=$?"
```

2.5 Your first plugin: a Git panel for web

dsh's sidebar is empty by default — install community plugin `dsh-better-sidebar` to feel "everything is a plugin" (mechanics in Ch.3; here just get it running):`

```
# 1. locate your web profile #    Windows: %USERPROFILE%\dsh\profiles\web #
macOS/Linux: ~/.dsh/profiles/web # 2. add one line to package.json dependencies #    "dsh-
better-sidebar": "link:C:\\path\\to\\DSH-better-sidebar" # 3. add to cordis.patch.yml #
- insert: #          - id: better-sidebar #          name: dsh-better-sidebar # 4. install &
restart cd ~/.dsh/profiles/web && pnpm install dsh web
```

After restart, the sidebar gains file manager / terminal / **Git panel** / browser tabs.

> The panel's fetch/pull/push buttons are a community PR (Ch.5) — **this is how the plugin ecosystem works**.

2.6 Config & paths

First run creates:

```
~/.dsh/ |— settings.yaml          # global settings (model, reasoning effort) |—
profiles/          # profile dirs |   |— web/ |   |— package.json          # plugin
deps + manifest |   |— cordis.patch.yml # your patch layer |— sessions/
# session data |— storages/          # persistence
```

``settings.yaml`` example:

```
agent-default-model:  model: deepseek-v4-flash  reasoningEffort: high
```

2.7 Command cheatsheet

Command	Purpose	
dsh web	start Web UI alias dsh --profile web	
dsh --profile headless "task"	one-shot task, print result, exit	
dsh plugin --profile <name> add <pkg>	add a plugin to a profile	
dsh --dump-config	print composed config tree	
dsh --profile tui	TUI mode plugin-provided; not built-in	
dsh --version	version	

2.8 Troubleshooting

Symptom	Cause & fix	
dsh: profile "tui" does not exist	tui needs a plugin dsh plugin --profile tui add <pkg>	
npx very slow	first-run package size; npm i -g helps	
browser can't reach 3080	port busy: netstat -ano \	findstr 3080 → kill PID
model not responding	check ~/.dsh/settings.yaml + API key	
plugin install 404	rc.1 dependency break: use ^0.1.0-rc.6 line Ch.3 pitfall #1	
behavior changed after upgrade	rc-stage breaking changes; check changelog	

****Next**:** [Chapter 3: Profiles & the Plugin System](./03-profiles.md)

[English](./03-profiles.en.md) | [中文](./03-profiles.md) · [[← Back](#)](../README.md)

Chapter 3: Profiles & the Plugin System

> ***Goal of this chapter:** Understand dsh's customizable skeleton, how profiles are organized, how plugins are mounted, and what the host/client dual halves are. ****This is the watershed moment from "user" to "developer."***

TL;DR (30-second version)

- 1. **Profile = a launchable form factor**: a `~/dsh/profiles/<name>/` directory composed of `package.json` (plugin manifest) + `cordis.patch.yml` (patch layer).
- 2. **Mounting a plugin takes two steps**: add a dependency in `package.json` + add a mount line in `cordis.patch.yml`, then `pnpm install` and restart.
- 3. **Host half runs in Node, client half runs in the browser**: one npm package carries both faces via `exports["."]` and `exports["./client"]`.
- 4. **To change behavior, look for extension points first, don't fork the core**: `agent/request` waterfall, `conversationEvents`, `ctx.slots.inject`, and the `settings` service are the four most common hooks.
- 5. **Three rc-stage pitfalls**: use the `^0.1.0-rc.6` dependency line, always `await next()`, and client tests need the dsh runtime.

3.1 Profile: A Launchable Configuration Stack

dsh uses **profiles** to represent "a launchable form factor." Two are built in; the rest are created via plugins:

Profile	Purpose	Command	
web	Web UI conversation + sidebar + tools	dsh web	
headless	One-shot CLI tasks	dsh --profile headless "task"	
tui plugin required	Terminal UI	dsh --profile tui not built in; requires a plugin	

A profile directory looks like this (`~/dsh/profiles/<name>/`):

```
<!-- [style] 目录树代码块统一补 text 语言标签 -->
profiles/web/ |— package.json      # Plugin dependencies + dsh.profile manifest (bundles
order) |— cordis.patch.yml      # Your patch layer: mount/override plugin declarations |—
cordis.yml      # (Generated) Composite configuration |— pnpm-workspace.yaml |—
node_modules/
```

****Loading order**** (from official docs): Built-in bundles (`dsh-base` → `dsh-web-app`) → profile's `cordis.patch.yml` → user-level `~/dsh/cordis.patch.yml` → `--patch` overlay.

3.2 Mounting a Plugin: Two Changes

Here's how to mount the community plugin `dsh-better-sidebar` (a real operation):

****① Add the dependency in `package.json`**** (using `link:` to point to local source, or an npm package name):

```
{  "dependencies": {    "dsh-better-sidebar": "link:C:\\path\\to\\DSH-better-sidebar"  } }
```

****② Add the mount line in `cordis.patch.yml`**:**

```
- insert:    - id: better-sidebar      name: dsh-better-sidebar
```

****③ Install and restart**:**

```
cd ~/.dsh/profiles/web && npm install dsh web # Takes effect after restart
```

> 💡 *Plugins are ****isolated per session**** by default (e.g. `better-sidebar`'s layout/tabs persist per session). Mounting is profile-level, but state is session-level.*

3.3 Host Half & Client Half: One Package, Two Faces

A plugin can carry two runtime halves simultaneously:

Half	Runs In	Responsibility	Example	
Host half	Node process	Tools, services, events, filesystem, processes	applyctx registers tools/services	
Client half	Browser web profile	UI, interaction, DOM	dsh.client declaration in package.json + src/client/	

How to declare the client half in `package.json`:

```
{  "dsh": {    "client": {      "inject": ["@deepseek-ai/dsh-client-runtime", "@deepseek-ai/dsh-client-locale"],      "platform": "web"    },    "exports": {      ".": { "types": "./lib/types/index.d.ts", "default": "./lib/index.js" },      "./client": { "types": "./lib/types/client/index.d.ts", "default": "./lib/client.js" }    }  }
```

- Host half: `exports["."]` → `apply(ctx)` (the cordis plugin body)
- Client half: `exports["./client"]` → `apply(ctx)` on the browser side
- `inject`: Declares required services (cordis dependency injection)

3.4 Extension Points: Look for Hooks First, Don't Fork the Core

Official principle (stated in CONTRIBUTING/AGENTS.md): *****Plugins, not loop changes: new behavior goes on documented extension points.***** The most common mistake newcomers make is modifying the core loop. The correct approach is to use extension points.

Confirmed commonly-used extension points (each explored hands-on in later chapters):

Extension Point	Location	Purpose	
agent/request waterfall	agent-loop	Modify config before every model request provider/model/reasoningEffort/tools. This is what the speed-up plugin example uses.	
agent/request-error	agent-loop	Intervene on request failure the official compaction plugin uses this for context-window overflow recovery	
conversationEvents.register	Client runtime	Subscribe to / inject conversation events tool/call, turn/start, etc.	
ctx.slots.inject	Client UI slots	Inject UI into interface slots e.g. turnTail to display artifact file lines	
settings service	dsh-settings	Register user-configurable namespaces settings page auto-renders them	
ctx.provide / ctx.get	cordis	Provide services across plugins	

3.5 Common Pitfalls (Battle-Tested)

Pitfall	Symptom	Fix	
rc.1 broken	pnpm install reports	Several packages from	

dependency chain	@deepseek-ai/dsh-type-meta@0.0.1-rc.1 404	the rc.1 era were never published; upgrade to the ^0.1.0-rc.6 line	
Plugin missing main	dsh: No "exports" main defined	The host plugin's package.json must expose a . entry "main": "src/index.ts" can be loaded directly by tsx	
Event handler forgets await next	Request loses provider/model and errors out	next in agent/request returns a Promise; you must await it before spreading	
Type error: 'agent/request' is not assignable to keyof Events	npm types don't re-export the official type augmentations	Use a relaxed signature ctx.on as unknown as ... at the boundary	
Client package depends on window.ModuleLoader	jsdom can't run client tests directly	Component-level tests need the dsh web runtime run in official CI	

Hands-on exercises

1. **Locate your profile**: find the `web` profile directory on your machine. Read `package.json` and `cordis.patch.yml`. Identify which plugins are currently mounted.
2. **Mount a plugin**: follow Section 3.2 to mount `dsh-better-sidebar` (or any community plugin). Verify it appears after restart.
3. **Trace the host/client split**: pick a plugin that has both halves. Find the `dsh.client` declaration in its `package.json` and the `apply(ctx)` function in its entry file.
4. **Find an extension point**: read the `agent/request` waterfall description above. In the dsh source, search for where this waterfall is called. What other waterfalls or hooks can you spot?
5. **Pitfall reproduction**: deliberately forget to `await next()` in a test plugin. Observe the error. Then fix it and confirm the config flows through.
6. **Think**: why does dsh separate "profile" from "plugin"? What would break if plugins were global instead of profile-scoped?

FAQ

- **Q: What's the difference between a profile and a plugin?** A profile is a **launchable form factor** (a directory with config + patch). A plugin is a **capability unit** mounted into a profile. You can think of a profile as a "container" and plugins as "modules inside it."
- **Q: Do I need to create a new profile for every project?** No. Most users stick with ``web`` for interactive work and ``headless`` for scripting. Custom profiles are for advanced use cases (TUI, desktop, bots).
- **Q: Why does ``pnpm install`` fail with 404 errors?** The rc.1 dependency chain is broken. Use the ``^0.1.0-rc.6`` line in ``package.json`` (see Section 3.5).
- **Q: Can I use a plugin without modifying ``cordis.patch.yml``?** No. Adding the dependency alone doesn't activate the plugin. You must declare it in the patch layer for Cordis to load it.
- **Q: Why can't I run client-half tests with jsdom?** The client half depends on ``window.__ModuleLoader__``, which is bootstrapped by the dsh web runtime. Component tests need the official CI environment.

Next chapter: [Chapter 4: Plugin Development, Hands-On](./04-plugin-dev.en.md) (planned) — Write your first host plugin from scratch.

[English](./04-plugin-dev.en.md) | [中文](./04-plugin-dev.md) · [[← Back](#)](../README.md)

Chapter 4: Plugin Development, Hands-On

> **Goal of this chapter:** Build a **real, working host plugin** from scratch, one that automatically adjusts reasoning effort via the ``agent/request`` extension point. This is a full breakdown of an example speed-up plugin. All code is runnable and testable.

TL;DR (30-second version)

1. **Core problem:** dsh re-thinks before every tool call; in a 50-step task, 90%+ of wall-clock time is thinking. Lowering the reasoning effort is the fastest speedup.
2. **Three-step architecture:** pure function (decision logic, zero deps, unit-testable) → plugin body (hook into the waterfall) → live verification (logs prove the injection happened).
3. **``agent/request`` waterfall:** fires before every model request; listener return values flow to the next listener, enabling "keep original config + override one field."
4. **``next()`` is a Promise; you must ``await`` it:** spreading without awaiting yields an empty object, dropping the provider/model and causing errors.
5. **Development discipline:** find the extension point first (90% of behaviors have official hooks), extract logic into pure functions, never skip live verification.

4.1 What We're Building

Problem: Before every tool call, the dsh model re-thinks from scratch (``reasoning_effort``). In a 50-step tool-chain task, "thinking" accounts for 90%+ of wall-clock time.

Solution: A host plugin that listens on the ``agent/request`` waterfall and downgrades ``reasoning_effort`` from ``high`` to ``low`` for simple turns, based on the most recent tool calls.

4.2 Project Skeleton

<!-- [style] 目录树代码块统一补 text 语言标签 -->

```
dsh-speed-plugin/ |— package.json          # Host plugin declaration |— tsconfig.json |—
src/ | |— effort-decision.ts # Pure function: decision logic (zero deps, unit-testable)
| |— index.ts              # apply(ctx): hooks into the extension point |— tests/ |—
effort-decision.spec.ts
```

Key fields in ``package.json``:

```
{  "name": "dsh-speed-plugin",  "type": "module",  "main": "src/index.ts",  "exports":
{    ".": { "types": "./src/index.ts", "default": "./src/index.ts" }  },
```

```
"peerDependencies": {    "@deepseek-ai/cordis": "^4.0.1",    "@deepseek-ai/dsh-agent": "^0.1.0-rc.6"  }
```

>  Always use the `^0.1.0-rc.6` dependency line. The rc.1 npm dependency chain is broken (see Chapter 3 pitfalls).

4.3 Pure Function: Decision Logic (Zero Dependencies, Unit-Testable)

`src/effort-decision.ts`:

```
export type EffortId = 'low' | 'high' | 'max' export interface ToolCallSample {  name: string    // Tool name, e.g. 'write', 'read', 'bash'  argsSize: number // Argument size (character count) } export interface EffortDecisionInput {  recentCalls: readonly ToolCallSample[]  selected: EffortId    // User's baseline effort  allowDowngrade: boolean  allowUpgrade: boolean } const SIMPLE_TOOL_RE = /^(fs|bash|terminal|read|write|grep|glob|edit|ls|cat|rm|cp|touch|mkdir|pwd)/i const HEAVY_ARGS = 800 export function decideEffort(input: EffortDecisionInput): EffortId {  const { recentCalls, selected, allowDowngrade, allowUpgrade } = input  if (recentCalls.length === 0) return selected  // Fresh prompt: keep baseline  const ratio = recentCalls.filter(c => SIMPLE_TOOL_RE.test(c.name) && c.argsSize < HEAVY_ARGS,  ).length / recentCalls.length  const heaviest = recentCalls.reduce((m, c) => Math.max(m, c.argsSize), 0)  if (ratio >= 0.75 && allowDowngrade) return 'low'  if (heaviest >= HEAVY_ARGS * 4 && allowUpgrade) return 'max'  if (ratio < 0.75) return allowUpgrade ? 'high' : selected  return selected }
```

****Why extract a pure function:**** The decision logic is decoupled from the dsh runtime. Unit tests have zero dependencies, run in milliseconds, and cover every branch. Live verification only needs to confirm "did the injection actually happen."

4.4 Plugin Body: Hooking into the `agent/request` Waterfall

`src/index.ts`:

```
import type { Context } from '@deepseek-ai/cordis' import { decideEffort, type ToolCallSample } from './effort-decision.ts' export interface SpeedPluginConfig {  enabled: boolean  allowDowngrade: boolean  allowUpgrade: boolean  baseline: 'low' | 'high' | 'max' } export const DEFAULT_CONFIG: SpeedPluginConfig = {  enabled: true,  allowDowngrade: true,  allowUpgrade: false,  baseline: 'high', } const WINDOW = 8 function recentToolCalls(agent: unknown): ToolCallSample[] {  const events = (agent as { session?: { events?: readonly unknown[] } }).session?.events ?? []  const out: ToolCallSample[] = []  for (let i = events.length - 1; i >= 0 && out.length < WINDOW; i--) {    const e = events[i] as { type?: string; data?: { name?: string; arguments?: unknown } } | undefined    if (e?.type !== 'tool/call') continue    out.push({      name: e.data?.name ?? 'tool',      argsSize: typeof e.data?.arguments === 'string' ? e.data.arguments.length : 0,    })  }  return out.reverse() } export function apply(ctx: Context, config: SpeedPluginConfig = DEFAULT_CONFIG): void {  if (!config.enabled) return  // Boundary adaptation: npm package doesn't re-export official event type augmentations,  // so we relax the signature here (see Chapter 3 pitfalls)  const on = ctx.on as unknown as (  event: string,  handler: (payload: Record<string, unknown>, next: () => unknown) => unknown | Promise<unknown>,  ) => void  on('agent/request', async (payload, next) => {    const seed = await next() as { reasoningEffort?: unknown }  //  Must await!    const calls = recentToolCalls(payload.agent)    const effort = decideEffort({      recentCalls: calls,      selected: config.baseline,      allowDowngrade: config.allowDowngrade,      allowUpgrade: config.allowUpgrade,    })    console.log(`[speed-plugin] calls=${JSON.stringify(calls)}`)    => reasoningEffort=${effort}`)    return { ...seed, reasoningEffort: effort }  }) }
```

****Three critical points**** (all learned the hard way):

1. ****`next()` is a Promise:**** ``await next()`` retrieves the current config. Spreading without awaiting yields an empty object, which drops the provider/model and causes errors.
2. ****Waterfall semantics:**** The listener's ****return value**** is passed to the next listener and ultimately to the request. Returning ``{...seed, reasoningEffort}`` means "keep the original config, but override the reasoning effort."
3. ****`agent/request` fires every step:**** The ``agent-loop``'s ``buildRequest`` runs through this waterfall at every step, so dynamic decisions naturally take effect per-step.

4.5 Testing

****Unit tests**** (pure function, zero dependencies):

```
import { describe, expect, it } from 'vitest' import { decideEffort } from '../src/effort-
decision.ts' it('downgrades to low for simple tool chains', () => { expect(decideEffort({
recentCalls: [{ name: 'write', argsSize: 40 }], selected: 'high', allowDowngrade: true,
allowUpgrade: true, })).toBe('low') }) // ... more branches: fresh prompt keeps
baseline / downgrade disabled / // oversized payload upgrades to max / mixed tools
upgrade to high
```

****Live verification**** (critical, proves "the injection actually happens"):

Mount the plugin (Chapter 3 method) → restart ``dsh web`` → send a file-creation task
→ watch the dsh process logs:

```
[speed-plugin] agent/request: calls=[] => reasoningEffort=high [speed-
plugin] agent/request: calls=[{"name":"write",...}] => reasoningEffort=low
```

First turn has no tool calls → stays at baseline ``high``. After detecting the ``write`` tool
→ next turn drops to ``low``. ****The injection pipeline works end to end.****

> Full runnable code: combine the code snippets in this chapter to run it.

4.6 Three Development Disciplines for Newcomers

1. ****Find the extension point first:**** 90% of behaviors you want to change have official hooks (``agent/request``, ``settings``, ``conversationEvents``, ``slots``). Don't fork the core.
2. ****Extract logic into pure functions:**** Decouple decision/computation logic from dsh. Unit tests run in milliseconds and cover every branch. Live verification only needs to confirm "the injection happened."
3. ****Never skip live verification:**** Unit tests prove the logic; live logs prove the wiring. Both must pass before you're done.

Hands-on exercises

1. **Read the pure function**: open `src/effort-decision.ts`. Trace the logic: what does `ratio >= 0.75` mean? When does it return `'low'` vs `'max'`?
2. **Write a unit test**: add a test case for "mixed tools with heavy arguments." What effort should it return? Run it.
3. **Break it on purpose**: remove the `await` before `next()` in `src/index.ts`. Run the plugin live. What error do you see? Fix it.
4. **Add a new rule**: extend `decideEffort` to always return `'max'` when the tool name is `'bash'` and the args exceed 2000 characters. Write a test for it.
5. **Live verification**: mount the plugin (Chapter 3 method), restart `dsh web`, send a file-creation task, and watch the logs. Confirm you see `reasoningEffort=low` after the first tool call.
6. **Think**: why extract the decision logic into a pure function instead of putting it directly in the `apply(ctx)` handler? What are the testing and maintenance implications?

FAQ

- **Q: Why not just set `reasoningEffort: low` globally?** Because complex tasks (planning, debugging) benefit from `high` or `max`. The plugin dynamically adjusts per step, giving you the best of both worlds.
- **Q: What if the plugin makes the wrong decision?** The plugin only downgrades when recent tool calls are simple (file ops, grep, etc.). If you're doing complex reasoning, the ratio drops and effort stays high. You can also disable downgrading via config.
- **Q: Can I use this plugin with other models?** The plugin hooks into dsh's `agent/request` waterfall, which is model-agnostic. It should work with any model dsh supports, though the speedup depends on the model's thinking behavior.
- **Q: Why is `next()` a Promise?** The waterfall is async: each listener may need to await the next one. Forgetting `await` breaks the chain and drops config.
- **Q: Where do I find more extension points?** Check `packages/AGENTS.md` in the official repo. The most common ones are listed in Chapter 3, Section 3.4.

Next chapter: [Chapter 5: Real-World Cases](./05-cases.en.md) (planned) — Git panel, HTML draft preview, speed-up plugin.

[English](./05-cases.en.md) | [中文](./05-cases.md) · [[← Back](#)](../README.md)

Chapter 5: What dsh Can Do for You — Application Scenarios

> *Goal: step out of the "how to build dsh" lens and look from the ****user's**** perspective — where dsh genuinely helps in daily work. Every scenario has a copy-paste task example with real measured timing.*

TL;DR (30-second version)

- 1. ****dsh is a general-purpose agent runtime**** — not just a coding assistant: data analysis, documentation, automation, and code comprehension all work.
- 2. ****Five core scenarios****: daily coding / data processing / documentation & knowledge / automation & ops / code understanding & refactoring.
- 3. ****One usage pattern****: describe the goal in one sentence + state the acceptance criteria; dsh orchestrates the toolchain itself.
- 4. ****Real speed****: simple tasks in seconds, complex multi-step tasks in 1–5 min (V4-Flash + high effort, measured).
- 5. ****Golden prompt rule****: tell dsh **what "done" looks like** ("run and verify") — more effective than describing the process.

5.1 Scenario overview

Scenario	Typical tasks	dsh tools	Measured	
Daily coding	write functions, fix bugs, add tests	write / read / bash / grep	5-bug fix + 49 tests: 94s	
Data processing	clean, analyze, visualize, report	read / write / bash / python	clean + visualize: 186s	
Docs & knowledge	API docs, tutorials, summaries, translation	read / write	API docs: 1m04s	
Automation & ops	batch, CI tasks, log analysis	bash / headless	headless scripting	
Code understanding	refactor, review, tech research	grep / glob / read / bash	refactor: 4m37s	

> *All timings from real cases in [Chapter 10](./10-complex-cases.en.md) (synthetic data); see [benchmark](./benchmark.md).*

5.1.1 High cache-hit rate: dsh's hidden cost weapon

****The most underrated advantage of the dsh ecosystem.**** DeepSeek's API ****context cache**** bills repeated/similar input tokens at a discounted rate. Measured in real dsh sessions: ****cache hit rate up to 97%****.

****Why dsh hits cache so often****:

Reason	Explanation	
Session continuity	multi-turn conversations repeat system prompt / skill catalog / history	
Stable tool schema	every step carries the same tool definitions	
Agent workload	multi-step tool chains repeatedly carry the same context	

****Cost impact (98% cache discount)****:

Charge	Miss	Cache hit	Discount	
Input Flash	\$0.14/M	\$0.0028/M	98%	
Input Pro	\$0.435/M	\$0.003625/M	99%+	

****Practical meaning****: in a long agent session, most input tokens hit cache — a 50-step tool chain with 2.4M input tokens costs far less than the miss-price estimate. For high-frequency / batch / long-session workloads this is an order-of-magnitude advantage.

> To raise hit rate: keep sessions alive, keep prompt prefixes stable (don't churn system prompt/skill catalog), batch within the same session/prefix.

5.2 Scenario 1: Daily coding (best entry point)

****Typical tasks****: write a function/module from a spec; fix bugs (locate → modify → verify); add unit tests.

```
dsh --profile headless "Review buggy_calculator.py: find all bugs, fix them, write complete tests, run and verify all pass" # measured: 5 bugs fixed + 49 tests pass (94s)
```

****Why dsh fits****: coding has the fullest toolchain — read files, write code, run commands, see results. Produced files are directly openable at the end of the conversation.

****Prompt tip****: always say "run and verify" — dsh closes the loop; "fix my bug" alone stops after editing.

5.3 Scenario 2: Data processing & analysis

```
dsh --profile headless "Analyze sales_data.json for quality issues
(missing/type/duplicate/outliers), write a cleaning script + visualization script, run and
verify, summarize your approach" # measured: 52→35 rows zeroed + chart + trade-off notes
(186s)
```

****Highlight**:** dsh doesn't just execute — it explains trade-offs ("median imputation creates a spike", "whether to drop outliers depends on business") — direct evidence of an agent that *thinks*.

5.4 Scenario 3: Documentation & knowledge work

```
dsh --profile headless "Generate complete Markdown API docs for user_module.py (8
functions): params, returns, exceptions, examples, edge cases" # measured: 423-line
ReadTheDocs-style doc (1m04s), proactively flags contract-vs-implementation gaps
```

****Highlight**:** defensive writing — notes calling traps (negative limit behavior, empty-argument behavior), flags where docs and code disagree.

5.5 Scenario 4: Automation & ops (headless is the ace)

```
# cron daily digest dsh --profile headless "Read today's git log in the workspace and write
a Chinese daily summary" > daily-report.md # batch: convert all .txt in docs/ to .md dsh --
profile headless "Convert all .txt files under docs/ to .md (preserve structure)"
```

****Why headless is the ace**:** full process-level tools + non-zero exit on failure + pipeable output — ****an agent you can put in CI****.

5.6 Scenario 5: Code understanding & refactoring

```
dsh --profile headless "Refactor legacy_orders.py to OOP with separated responsibilities,
keep behavior identical, write tests verifying output parity" # measured: 17/17 tests pass
incl. script-level artifact comparison (4m37s)
```

****Highlight**:** engineering awareness — merges duplicated logic into one entry point, keeps backward compatibility, ****notices a sandbox tempfile-permission issue and switches approach****.

5.7 Three prompt rules that make scenarios work

1. ****State acceptance criteria, not process**:** "run and verify" > "write a script" — dsh figures out the steps.
2. ****Give context**:** one sentence of background (where files are, what the data looks like, who the audience is).
3. ****One task at a time**:** split big goals into rounds of small closed loops.

5.8 Your first real task (15 minutes)

1. Make a test dir with a small file (code / data / notes).
2. Run: ``dsh --profile headless "Summarize this file and list 3 improvement points"``.
3. Run: ``dsh --profile headless "Write a test / generate docs for this file"``.
4. Compare output — did it give a real artifact, or a vague answer? Were your acceptance criteria clear enough?

5.9 Industry view: how dsh opens in different domains

> Same toolchain, different industries. Generic scenarios only (no real business data); regulated fields (medical/legal) require human review.

Finance & research

``dsh --profile headless "Read this financial data and generate a summary with YoY/MoM changes"`` — data pipeline + high cache hit make long-session reports cheap.

Education & learning

``dsh --profile headless "Turn this chapter into 10 self-test questions with explanations"`` — document generation + structuring are strong; batch question generation.

Sales & support

``dsh --profile headless "Group 100 customer feedback entries by issue type and suggest improvements"`` — batch text processing + classification.

Operations & content

``dsh --profile headless "Read this week's data files and generate an ops weekly report (with chart code)"`` — data + docs + automation combined.

Research & review

``dsh --profile headless "Compare the methods and conclusions of these 3 papers, output a comparison table"`` — long-doc processing + structured output (1M context).

⚠ Compliance note

High-risk domains (medical diagnosis, legal advice, financial advice): dsh output needs human review — tools are sandboxed, but **content responsibility is on the user**.

Hands-on exercises

1. **Scenario mapping**: classify 3 of your daily tasks into the five scenarios.
2. **Prompt comparison**: write two prompt versions for the same task — one describing process, one stating acceptance criteria — run both and compare.
3. **Headless scripting**: turn "daily git digest" into one cron/bash command; verify output.
4. **Multi-step task**: design a cross-scenario task (read code → generate docs → output report) and watch dsh orchestrate.
5. **Think**: which scenarios is dsh NOT suitable for? (Hint: real-time interaction, sensitive operations needing human approval.)

FAQ

- **Q: Is dsh just a coding assistant?** No — its toolchain (read/write/run/search) makes it general; data/docs/automation all verified.
- **Q: How long do complex tasks take?** 1–5 min measured (V4-Flash); lower reasoning effort for speed (Ch.6).
- **Q: What if a task fails?** Trace the tool calls in output; usually the acceptance criteria were unclear — rewrite and rerun.
- **Q: Can it run automated?** Yes — headless + non-zero exit + pipeable output works in cron/CI.
- **Q: Is it safe?** Tool execution has sandbox + approval layers; sensitive ops should be human-confirmed (Ch.8).
- **Q: How is this different from using Claude Code directly?** Same-model capability is close, but dsh lets you customize toolchains/UI/plugins (Ch.1 matrix).

Next: [Chapter 6: Advanced & Performance Tuning](./06-advanced.en.md)

[English](./06-advanced.en.md) | [中文](./06-advanced.md) · [[← Back](#)](../README.md)

Chapter 6: Advanced & Performance Tuning

> **Goal of this chapter:** Go from "it runs" to "it runs well" — reasoning effort strategy, tool-call latency analysis, and a real-world pitfall checklist.

TL;DR (30-second version)

- Where the time goes:** model thinking takes ~90%, tool execution <1%, network/rendering ~10%. Optimizing thinking time is the highest-leverage speedup.
- Three-tier strategy:** `low` (simple/batch turns), `high` (everyday default), `max` (complex reasoning/debug). Choose manually via the UI or let a plugin adjust automatically.
- Latency visualization:** the stats line at the bottom of the Web UI (LLM Xs / Tool calls Ys) is the fastest way to locate bottlenecks.
- Seven real pitfalls:** rc.1 dependency break, missing plugin `main`, un-awaited `next()`, unrecognized event types, client tests that won't run, cache-hit misattribution, port conflicts.
- Evaluation three questions:** who tested it, which harness, how strict is the verifier. Always read benchmarks with context.

6.1 Performance Model: Where Does dsh Spend Its Time

Measured latency breakdown for a "create a file" task:

Phase	Share	Notes	
Model thinking Think	~90%	Re-thinks before every tool call — the overwhelming majority	
Tool execution	<1%	File writes and similar, millisecond-level	
Network / rendering	~10%	API round-trip + UI updates	

Implications:

- Simple tasks → optimize thinking time (lower reasoning effort)
- Long tool-chain tasks → save thinking time at every step; cumulative gains are largest
- Slow tools themselves (search / large files) → optimize the tool implementation, not the effort level

6.2 reasoning_effort Strategy (Official Three Tiers)

Tier	Recommended Use	
low	Simple / deterministic turns: file operations, batch jobs, cheap steps in a tool chain	
high	Everyday agent tasks default	
max	Complex reasoning, long-chain planning, debugging	

- Manual:** The "Reasoning Level" selector in the UI, or `reasoningEffort` in `~/.dsh/settings.yaml`.
- Automatic:** A plugin dynamically downgrades per tool turn (example speed-up plugin, Chapter 4).

6.3 Tool-Call Latency Visualization

dsh's session stats line (at the bottom of the Web UI) shows: ``N turns · M steps | LLM Xs · Tool calls Ys | Avg first token ...`` — this is the fastest way to locate bottlenecks.

Advanced: A host plugin can listen to tool events for per-tool timing (the example speed-up plugin already includes a host-side logging version), surfacing "which tool is the slowest."

6.4 Pitfall Checklist (Battle-Tested, with Fixes)

#	Pitfall	Symptom	Fix	
1	rc.1 broken dependency chain	pnpm install 404 dsh-type-meta etc. were never published	Use the ^0.1.0-rc.6 dependency line	
2	Plugin missing main	No "exports" main defined	Expose a . entry; "main": "src/index.ts" can be loaded by tsx	
3	next not awaited	Provider/model lost, error thrown	next in agent/request returns a Promise; must await	
4	Event type not recognized	'agent/request' is not assignable to	npm doesn't re-export type	

		keyof Events	augmentations; relax the signature at the boundary	
5	Client tests won't run	jsdom reports window.ModuleLoader undefined	Client artifacts depend on dsh's bootstrap mechanism; run component tests in official CI	
6	Misattributing "suddenly faster" on simple tasks	1s vs 110s difference misattributed	DeepSeek context cache hits also speed things up — A/B tests must use a fresh prompt	
7	Port conflict	dsh web won't start	netstat -ano	findstr 3080 to find the PID, then kill it

6.5 Evaluation Perspective: Official Scorecards vs. Independent Benchmarks

(With the 0813 official release in mind) When reading agent model scorecards, ask three questions:

1. **Who tested it?** Official self-tests (their own harness) vs. independent third parties (AA, etc.)
2. **Which harness?** Different frameworks yield very different scores (official Terminal-Bench 87.9 vs. AA independent 79)
3. **How strict is the verifier?** Lenient verifier (SWE-bench Verified has 8.5% false positives) vs. strict (DeepSWE 0.3%)

dsh's `agent/request` waterfall makes model benchmarks reproducible. That's an engineering advantage over closed-source products.

Hands-on exercises

1. **Measure your own task**: run a simple file-creation task in `dsh web`. Check the stats line at the bottom. What's the LLM time vs tool-call time? Does it match the 90% / <1% split?

2. **Effort comparison**: run the same task twice, once with ``reasoningEffort: low`` and once with ``high``. Compare wall-clock time and output quality. When is ``low`` good enough?
3. **Pitfall hunt**: open the pitfall checklist (Section 6.4). For each of the 7 pitfalls, try to reproduce it (or find a log/example where it occurred). Write down the fix.
4. **Plugin timing**: if you've built the example speed-up plugin (Chapter 4), check its logs. How much time does each step save? Calculate the cumulative gain over a 50-step task.
5. **Benchmark skepticism**: find an official DeepSeek benchmark scorecard. Ask the three evaluation questions (Section 6.5). What harness did they use? How strict was the verifier?
6. **Think**: why does dsh re-think before every tool call? What would happen if it cached the reasoning across steps? What are the trade-offs?

FAQ

- **Q: Why is model thinking 90% of the time?** Because dsh uses a reasoning model that generates a chain-of-thought before every action. In a multi-step task, each step triggers a new reasoning pass. This is by design for quality, but it's slow.
- **Q: Can I just set ``reasoningEffort: low`` forever?** For simple, deterministic tasks (file ops, batch jobs), yes. For complex reasoning, planning, or debugging, you'll get better results with ``high`` or ``max``. The plugin approach (Chapter 4) gives you dynamic adjustment.
- **Q: Why did my task suddenly run 100x faster?** Likely a context-cache hit. DeepSeek's API caches repeated input tokens at a 98% discount. If your prompt/history is similar to a previous request, most tokens hit cache. This is a good thing, but it makes A/B testing tricky (use fresh prompts).
- **Q: What's the "overflow agent" in compaction?** When a conversation exceeds the context window and compression fails, dsh routes to a fallback "overflow agent" that tries to salvage the task. It's a last-resort mechanism.
- **Q: How do I read the stats line?** ``N turns · M steps | LLM Xs · Tool calls Ys | Avg first token ...`` — LLM time is total model thinking, tool-call time is total tool execution. If LLM time dominates, lower the effort. If tool time dominates, optimize the tool.

- **Q: Are official benchmarks reliable?** They're a starting point, but always ask: who tested it, which harness, how strict is the verifier. Independent benchmarks with strict verifiers are more trustworthy.

Next chapter: [Chapter 7: Ecosystem & Resources](./07-ecosystem.en.md) (planned).

[English](./07-ecosystem.en.md) | [中文](./07-ecosystem.md) · [[← Back](#)](../README.md)

Chapter 7: Ecosystem & Resources

> **Goal of this chapter:** Give you a map for "joining the dsh ecosystem" — official entry points, the state of the community, and how to participate.

TL;DR (30-second version)

- 1. **Official entry points:** GitHub repo (source + issues), API docs, Discord, Discussions. The current contribution entry is primarily via Discussions.
- 2. **External PRs are not accepted yet** (as of 2026-08-13), but the official team encourages dsh-plugin ecosystem projects, tutorials, and blog posts.
- 3. **Plugin ecosystem snapshot:** `DSH-better-sidebar` (most complete community plugin), an example speed-up plugin (Chapter 4 full breakdown), this handbook. Search `topic:dsh-plugin` to discover more.
- 4. **Beginner participation path:** use it → make small improvements (PRs to community plugins) → publish your own plugin → write content.
- 5. **Ecosystem day-zero = first-mover advantage:** every early ecosystem has a "first crab" bonus. Now is the time to get in.

7.1 Official Entry Points

Resource	URL	Purpose	
Official repository	https://github.com/deepseek-ai/deepseek-harness	Source code, architecture docs, issues	
API documentation	https://api-docs.deepseek.com	Models, pricing, API guides	
Discord	Link in official README	Community discussion	
Discussions	Official repo Discussions	Proposals / help requests currently the recommended contribution entry point	

Note: The official CONTRIBUTING guide states "external PRs are not accepted at this time" (as of 2026-08-13). However, the following are encouraged:

- Submit suggestions in Discussions (the team evaluates them)
- **Build dsh-plugin ecosystem projects** (an officially endorsed contribution path)
- Write tutorials and blog posts

7.2 Plugin Ecosystem (Snapshot as of 2026-08-13)

Project	Focus	Status	
DSH-better-sidebar	File management / terminal / Git / browser sidebar	Most complete community plugin	
Example speed-up plugin	Tool-call speed-up automatic reasoning effort adjustment	Teaching example Chapter 4	
dsh-handbook this handbook	Beginner tutorial	Ecosystem documentation	

****Discovering plugins:**** Search GitHub for ``topic:dsh-plugin``.

****Publishing a plugin:**** Add the ``dsh-plugin`` topic to your repo + publish on npm.

7.3 How to Join the Ecosystem (Beginner Path)

1. ****Use it:**** ``dsh web`` + install two community plugins, integrate into your daily workflow
2. ****Small improvements:**** Submit PRs to community plugins (read the three cases in Chapter 5 for a complete PR paradigm)
3. ****Publish a plugin:**** Start from the minimal host plugin in Chapter 4, tag it with ``dsh-plugin``
4. ****Write content:**** Tutorials / reviews / pitfall guides (officially encouraged); cross-reference this handbook

7.4 Recommended Reading Path

Goal	Path	
Quick start	Chapter 2 → install better-sidebar → use daily	
Plugin development	Chapter 3 → Chapter 4 → follow the Chapter 5 cases	
Performance tuning	Chapter 6 → Chapter 4 example source	
Deep customization	Official AGENTS.md architecture → docs/architecture.md → packages/ source code	

Closing Words

dsh was open-sourced on 2026-08-13. **Every day of the ecosystem is "early days."** This handbook will keep updating as dsh evolves. If a command in some chapter stops working, it's most likely due to rc version iteration. Always defer to the official changelog.

Hands-on exercises

1. **Explore the official repo**: go to <https://github.com/deepseek-ai/deepseek-harness>. Read the README, then open ``packages/AGENTS.md``. What architecture insights can you find?
2. **Join the community**: join the Discord or browse the Discussions tab. What topics are people talking about? What questions are unanswered?
3. **Install a community plugin**: if you haven't already, install ``dsh-better-sidebar`` or build the example speed-up plugin from Chapter 4. Use it for a day. What works well? What could be improved?
4. **Submit a suggestion**: go to the official Discussions tab. Write a proposal for a feature or improvement. Be specific: what problem does it solve? How would it work?
5. **Publish a plugin**: start from the minimal host plugin in Chapter 4. Add a feature (e.g. a new tool, a UI tweak). Tag it with ``dsh-plugin`` and publish to npm.
6. **Think**: why does the official team encourage ecosystem projects instead of accepting PRs? What are the advantages and risks of this approach?

FAQ

- **Q: Can I contribute code to the official repo?** Not via PR at the moment (2026-08-13). The official team encourages ecosystem projects (plugins, tutorials, tools) instead. Check the CONTRIBUTING guide for updates.
- **Q: How do I discover community plugins?** Search GitHub for ``topic:dsh-plugin``. The ecosystem is growing fast, so new plugins appear regularly.
- **Q: What's the best way to get help?** Join the Discord or post in the official Discussions tab. The community is active and responsive.
- **Q: Do I need to publish my plugin to npm?** No, but it's recommended. If you want others to use it easily, publish to npm and tag it with ``dsh-plugin``.

- **Q: Is it too late to join the ecosystem?** No. The ecosystem was open-sourced on 2026-08-13. Every day is "early days." First-movers have an advantage, but the ecosystem is still small enough that quality contributions stand out.

Go claim your spot in this brand-new ecosystem. 

[English](./08-tools-context.en.md) | [中文](./08-tools-context.md) · [[← Back](#)](../README.md)

Chapter 8: Tools & Context System

> **Goal of this chapter:** Understand dsh's "capability engine" — which tools the model can call, how context is fed to the model, and how long conversations are handled. **This is the key chapter for going from "it runs" to "I understand how it runs."**

TL;DR (30-second version)

- 1. **60+ official capability packages:** tools (fs/shell/web/skill/todo), context (context/compaction), session, subagent, MCP, workflow, safety, model, UI. Everything is a plugin.
- 2. **Built-in tool names are short verbs:**
``read`/`write`/`grep`/`glob`/`edit`/`bash`/`todo`/`skill``. When writing prompts, just say "read the file" or "search" and it works.
- 3. **Tool returns carry `locations` → artifact tracking:** the model can see which files it changed, and the UI lets you open them directly (artifact chips at the end of the conversation).
- 4. **Context = system prompt + skill catalog + conversation history + tool results:** layered injection, with tool schemas carried in every request.
- 5. **Long conversations are auto-compressed (compaction):** detect overflow → prune history → optional summarization → fallback to overflow agent. Important info should be written into the prompt manually.

8.1 Official Capability Map (60+ Packages at a Glance)

All of dsh's capabilities are provided as packages (``packages/<group>/<name>``). The ones newcomers need to know first:

Capability Domain	Official Packages	Purpose	
Tools	fs/tool-fs, fs/tool-fs-search, fs/tool-str-replace-editor, shell/tool-bash, web, skill, todo	Callable tools for files, terminal, web, skills, todos, etc.	
Context	context/, compaction/	Request context assembly, long-conversation compression	
Session	session/	Persistence, titles, telemetry	

Subagent	subagent/	Delegate sub-tasks	
MCP	mcp/	MCP client external tool servers	
Workflow	workflow/	Multi-step workflow orchestration	
Safety	sandbox/, guard/, interaction/	Sandboxing, loop hygiene, permissions/approvals	
Model LLM	llm/, llm-deepseek, llm-retry	Model integration, retries	
Skill	skill/	Skill provider registry	
Client UI	client/ ui-conversation, ui-tool, ...	Web UI components	

> Full list: see ``packages/AGENTS.md`` in the official repository.

8.2 Built-In Tools (Observed in Practice)

The actual tool names the model can call are **short verbs** (observed from dsh web sessions and agent/request logs):

Tool Name	Purpose	Notes	
read	Read files	fs capability	
write	Write files	fs capability	
grep	Content search	fs-search	
glob	File pattern matching	fs-search	
edit / strreplaceeditor	Precise editing	Tool results include locations used for artifact tracking	
bash / pwsh	Execute commands	Sandbox isolation	
todo	Todo management	Long-task planning	
skill	Skill invocation	Injected via skill-catalog	

Tool results and artifact tracking (important concept): Tool return values carry ``locations`` (file paths). dsh uses these to build "artifact file lines" — the artifact chips at the end of a conversation come from here. **The model can see which files were changed, and the UI lets you open them directly.**

8.3 How Context Is Fed to the Model

A single model request's context = system prompt + skill catalog + conversation history + tool results. This is visible in session logs:

<!-- [style] 示意图代码块统一补 text 语言标签 -->

```
Context injection @deepseek-ai/dsh-system-prompt ← Official system prompt Context
injection skill-catalog ← Skill catalog
```

dsh's context mechanism:

- **Layered system prompts:** Official plugins register prompt sections via ``systemPrompt.section()`` (e.g. ui-deliverables registers "artifact file reference" guidance)
- **Skill catalog injection:** The list of available skills enters the context; the model calls them as needed
- **Tool schemas:** Every request carries tool definitions

8.4 Long Conversations: Compaction

Long conversations blow up the context window. dsh's ``compaction`` plugins (e.g. ``compaction-basic``) handle:

- Detecting context overflow (``CONTEXT_WINDOW_EXCEEDED`` via ``agent/request-error``)
- Compressing history (model-agnostic pruning + optional summarization)
- Routing to an "overflow agent" on failure

*> For newcomers: **Just know that "long conversations are auto-compressed."** The details are an advanced topic. In production, note that compression loses detail — important context should be written into the prompt manually.*

8.5 Permissions & Security Model (Awareness Level)

- **Access modes:** The UI shows modes like "Workspace Write" (permission presets)
- **Interactive approvals:** ``interaction/*`` provides permission/approval capabilities (dangerous operations can require confirmation)
- **Sandboxing:** ``sandbox/*`` isolates command execution (e.g. the pwsh sandbox has ACL constraints — temp directory permission issues have been observed in practice)
- **Tool timeouts:** ``guard/*`` provides loop hygiene and tool timeouts

*> Deep security configuration is beyond the scope of this handbook. The key takeaway: **dsh's tool execution has isolation and approval layers by default.** It's not bare execution.*

8.6 Three Things Newcomers Should Remember Most

1. **Tool names are short verbs** (read/write/grep/glob/edit/bash) — when writing prompts or plugins, just say "read the file" or "search" and it works
2. **Tool returns include `locations` → artifact tracking** — files the model changed appear in the conversation's artifact area
3. **Long conversations are auto-compressed** — no need to manually clear history (but important info should go into the prompt)

Hands-on exercises

1. **Tool inventory**: open a `dsh web` session. Ask the model: "List all the tools you can call." Compare the list with Section 8.2. Are there any surprises?
2. **Artifact tracking**: give dsh a task that modifies multiple files (e.g. "Create a Python project with a main script, a test file, and a README"). After the task, check the artifact chips at the end of the conversation. Can you open each file?
3. **Context inspection**: open the session log (top-right in `dsh web`). Look for "Context injection" lines. What sections are injected? How much of the context is system prompt vs conversation history?
4. **Long conversation test**: have a 20+ turn conversation with dsh. At what point does compaction kick in? Check the logs for compaction events. Does the model still remember early context?
5. **Permission modes**: in the Web UI, switch between different access modes (e.g. "Workspace Write"). What changes? What operations are restricted?
6. **Think**: why are tool names short verbs instead of descriptive names? How does this affect the model's ability to use them correctly?

FAQ

- **Q**: What's the difference between `edit` and `write`? **A**: `write` creates or overwrites a file. `edit` (or `str\_replace\_editor`) makes precise replacements within a file. Use `edit` for small changes, `write` for new files or major rewrites.
- **Q**: Why does the model sometimes use `bash` instead of `read`? **A**: The model chooses tools based on the task. If it needs to run a command (e.g. `cat file.txt`),

it uses ``bash``. If it just needs to read the file content, it uses ``read``. Both work, but ``read`` is more efficient.

- **Q: What happens when the context window fills up?** dsh's compaction plugins detect the overflow, prune the history, and optionally summarize. You don't need to manually clear the conversation, but important context may be lost. For critical info, write it into the prompt.
- **Q: Can I add custom tools?** Yes, via a host plugin. Use the ``tools`` capability to register new tools. The model will see them in the tool schema and can call them.
- **Q: What's a "skill" and how is it different from a tool?** A skill is a higher-level capability injected via the skill catalog. The model calls skills via the ``skill`` tool. Skills are typically more complex than tools (e.g. "review this code" vs "read this file").
- **Q: Is sandboxing enabled by default?** Yes. Tool execution has isolation and approval layers. Dangerous operations (e.g. ``bash`` commands) may require confirmation. You can configure the sandbox via the ``sandbox/*`` and ``interaction/*`` packages.

Next chapter: [Chapter 9: MCP, Subagents & Workflows](./09-mcp-subagent-workflow.en.md)

[English](./09-mcp-subagent-workflow.en.md) | [中文](./09-mcp-subagent-workflow.md) · [[← Back](#)](./README.md)

Chapter 9: MCP, Subagents & Workflows

> **Goal of this chapter:** Understand dsh's three extension capabilities — MCP (connecting external tools), subagents (parallel tasks), and workflows (multi-step orchestration). **These are the switches that upgrade dsh from a "single agent" to an "agent system."**

>  This chapter is based on official architecture docs (`packages/AGENTS.md`) and package structure. Some feature examples are marked "pending live testing" — rc versions iterate quickly, so defer to the official changelog for exact usage.

TL;DR (30-second version)

1. **MCP = plug external tools into the agent**: connect to databases, browsers, internal systems via the MCP protocol. Get the built-in tools working first, then add MCP when you have a clear gap.
2. **Subagents = work in parallel**: delegate tasks to sub-agents for parallel execution (large-repo research, long-task decomposition, independent verification). Master single-agent workflow first.
3. **Workflows = deterministic process orchestration**: unlike "multi-turn conversation" (model freestyles), workflows define steps as a flow executed in order or by condition. Suited for fetch → clean → report → validate.
4. **Combined = agent system**: parent agent plans → subagents research in parallel → tool chain + MCP → workflow aggregates → artifacts.
5. **Beginner path, four stages**: single agent + built-in tools → + MCP → + subagents → + workflows. Stack gradually, don't skip levels.

9.1 MCP: Connecting to the External Tool Ecosystem

MCP (Model Context Protocol) is an open protocol for "plugging external tools into an agent." dsh provides an `mcp` client package (@deepseek-ai/dsh-mcp-client).`

What it does: Connect to any MCP server via MCP (databases, browsers, charts, internal company systems, etc.). For example, the community's `dsh-plugin-cost-tracker` (token tracking) is an MCP/plugin hybrid.`

Positioning for newcomers: MCP is an "ecosystem extension." Only reach for it once you have a clear tool gap inside dsh. Get the built-in tools working first, then add MCP.

9.2 Subagents: Working in Parallel

What they are: Delegate tasks to sub-agents for parallel execution (``packages/subagent/*``).

Typical scenarios:

- Multi-module research in a large repo (one subagent per module)
- Long-task decomposition (parent agent plans, subagents execute)
- Independent verification (subagents cross-check each other)

For newcomers: Subagents are "advanced weaponry." Master the single-agent workflow first, understand task decomposition, then move to subagents. A common anti-pattern is "spawning a subagent for everything," which just adds coordination overhead.

9.3 Workflows: Multi-Step Orchestration

What they are: ``packages/workflow/*`` provides multi-step workflow orchestration (worker-thread provider + tool Consumer).

How they differ from "multi-turn conversation": A normal conversation is "the model freestyles." A workflow is "steps are defined as a flow, executed in order or by condition." This suits **deterministic processes** (e.g. fetch data → clean → generate report → validate).

Official examples (``packages/examples/``): Runnable `cordis.yml` workflow samples are available.

9.4 Putting It Together: What an "Agent System" Looks Like

<!-- [style] 流程图代码块统一补 text 语言标签 -->

```

Your prompt (goal)  ↓ Parent Agent (planning)  |— Subagent A: Research module A
(parallel)         |— Subagent B: Research module B (parallel)  |— Tool chain: read/grep/bash
+ MCP (database)    ↓ Workflow (e.g. validate → aggregate → output report)  ↓ Artifacts
(openable / trackable)

```

Suggested beginner path:

1. **Stage 1:** Single agent + built-in tools (Chapters 2–5)
2. **Stage 2:** + MCP (connect external tools)
3. **Stage 3:** + Subagents (parallelism)
4. **Stage 4:** + Workflows (deterministic flows)

9.5 Community Ecosystem Examples (Snapshot as of 2026-08-13)

Community projects already appearing in the official Discussions:

- `dsh-plugin-cost-tracker` — real-time token cost tracking (plugin/MCP hybrid)
- Plugin development / speed-up tools (e.g. the example speed-up plugin in Chapter 4 of this handbook)

The ecosystem is growing fast. ****Now is a great time to start building.****

Hands-on exercises

1. ****MCP exploration****: search GitHub for MCP servers (e.g. database, browser, chart). Pick one. Read its documentation. What would it add to dsh?
2. ****Subagent design****: think of a task that could benefit from parallelism (e.g. "research 5 modules in a large repo"). How would you decompose it into subagent tasks? What would the parent agent do?
3. ****Workflow sketch****: design a deterministic workflow for a real task (e.g. "daily data pipeline: fetch → clean → analyze → report"). What are the steps? What conditions or branches exist?
4. ****Agent system diagram****: draw a diagram of an "agent system" that combines all four capabilities (single agent + MCP + subagents + workflow). What does each component do?
5. ****Beginner path check****: which stage are you at? If you're at stage 1 (single agent), master it first. If you're at stage 2 (MCP), what external tool are you connecting? What gap does it fill?
6. ****Think****: why does the guide say "don't skip levels"? What goes wrong if you spawn subagents before mastering single-agent workflow?

FAQ

- ****Q: When should I add MCP?**** Only when you have a clear tool gap. If dsh's built-in tools (read/write/bash/grep) can't do what you need (e.g. query a database, interact with a browser), then MCP is the answer.
- ****Q: Do subagents speed up every task?**** No. Subagents add coordination overhead. They're useful for parallelizable tasks (research, decomposition, verification). For sequential tasks, they just add complexity.

- **Q: What's the difference between a workflow and a multi-turn conversation?**
A conversation is "the model freestyles." A workflow is "steps are defined as a flow." Workflows are deterministic; conversations are adaptive. Use workflows for repeatable processes.
- **Q: Can I combine subagents and workflows?** Yes. A workflow can spawn subagents for parallel steps. This is the "agent system" pattern: parent plans, subagents execute, workflow orchestrates.
- **Q: Are subagents and workflows stable in rc?** They're part of the official architecture, but rc versions iterate quickly. Check the changelog and ``packages/AGENTS.md`` for the latest status.
- **Q: Where do I find MCP servers?** Search GitHub for "MCP server" or check the official MCP documentation. The ecosystem is growing fast.

Appendix A: [Glossary & Command Quick Reference](./appendix-glossary.md)
[English](./10-complex-cases.en.md) | [中文](./10-complex-cases.md) · [[← Back](#)](../README.md)

Chapter 10: Complex Real-World Cases (Run Live in dsh)

> **Goal of this chapter:** Demonstrate dsh's actual capabilities, output quality, and timing on multi-step tool chains through two **complex tasks completed live in dsh**.
> **Privacy notice:** Both cases use **synthetic data / self-authored code** only. No real business data or sensitive information is involved.

TL;DR (30-second version)

- Case A (data quality analysis, 186 seconds)**: 52 rows of dirty data → analysis → cleaning script → visualization → run verification. 52 → 35 rows, all issues resolved.
- Case B (5-bug fix + 49 tests, 94 seconds)**: calculator module with 5 planted bugs → all fixed → 49 tests pass, covering edge cases, precision, and exceptions.
- dsh's profile**: auto-orchestrated multi-step tool chains, shows judgment (proactively explains data trade-offs), artifact tracking, controllable timing (94-186s).
- Key insight**: write acceptance criteria clearly in the prompt (e.g. "run and verify"), and dsh will close the loop. For complex tasks, specify "what to do + how to verify."
- Failure recovery**: not triggered in these cases. Production recommendation: pair with guard timeouts + human approval.

Case A: Data Quality Analysis + Cleaning + Visualization (186 seconds)

Task

Give dsh a synthetic JSON file with dirty data (52 rows: missing values, type errors, duplicate rows, outlier value 99999), and ask it to:

- Analyze data quality issues
- Write a cleaning script ``clean.py``
- Write a visualization script ``visualize.py`` (generating ``chart.png``)
- Run both scripts to verify

How dsh Did It (Tool Chain)

<!-- [style] 工具链示意代码块统一补 text 语言标签 -->

```
read(sales_data.json) → write(clean.py) → bash(python clean.py) → write(visualize.py) → bash(python visualize.py) → read(output) → summary
```

Output & Verification

Artifact	Result	
clean.py	✔ Runs successfully, 52 → 35 rows, all issues resolved	
visualize.py + chart.png	✔ Valid PNG generated 1950×825	
Cleaning strategy	Median imputation for missing amounts, outlier removal, deduplication, type normalization	

![dsh-generated sales visualization (Case A artifact)](/assets/case-chart.png)

Output Highlights (Shows Judgment, Not Just Mechanical Execution)

> - Median imputation creates a 510 spike in the histogram — an inherent side effect of imputation, worth noting

> - If 99999 represents a genuine large order, consider IQR/quantile methods instead of blanket deletion

> - Post-cleaning stats: N=35, mean 489.43, median 510, std dev 181.73

****dsh didn't just execute the task — it proactively explained the trade-offs and assumptions in data processing.**** This is evidence of "thinking" at the agent engineering layer.

Case B: 5-Bug Code Fix + 49 Tests (94 seconds)

Task

Give dsh a Python calculator module ****deliberately seeded with 5 bugs****, and ask it to: find all bugs → fix them → write comprehensive unit tests → run and verify.

The Planted Bugs

- 1. `add` mistakenly written as `a - b`
- 2. `divide` silently returns 0 on division by zero (should throw)
- 3. `factorial` silently returns -1 for negative input (should throw)
- 4. `factorial`'s `range(1, n)` misses multiplying by n
- 5. `format\_money` doesn't format to two decimal places

dsh's Fixes + Tests (Verification: 49 passed)

Full test file: [case-test-calculator.py](./assets/case-test-calculator.py) — 49 test cases covering:

Function	Coverage	
add	Positive / negative / zero / float / precision / commutativity	
subtract	Negative results / zero boundary / float	
multiply	Sign combinations / multiply by zero / float	
divide	Integer division / signs / float / division by zero must throw ZeroDivisionError	
factorial	0!/1! boundary / known values / negative must throw ValueError	
formatmoney	Zero-padding / rounding / 0.1+0.2 rounding / negative	

python -m pytest test_calculator.py -v # 49 passed in 0.37s

Observations: dsh's Performance

- **Full closed loop:** Read → fix → write tests → run tests → summarize, with no human intervention
- **Fix quality:** All 5 bugs fixed, with docstring explanations added
- **Test design:** Covers edge cases (division by zero / negatives / precision), not just "does it run"

Case Summary: dsh's Profile on Complex Tasks

Dimension	Performance	
Multi-step tool chain	✅ Auto-orchestrated read → write → run → verify → summarize	
Complex task timing	94s–186s same model, same gateway; slower than benchmark simple tasks but manageable	
Output quality	✅ Shows judgment data trade-off explanations, test edge-case	

	design	
Artifact tracking	 Generated files can be opened directly at the end of the conversation	
Failure recovery	Not triggered in these cases; production recommendation: pair with guard timeouts + human approval	

Takeaway for newcomers: dsh excels at tasks involving "multiple files, multiple steps, and verification" — which is also its primary value proposition as an agent runtime. For complex tasks, the recommendation is to **write acceptance criteria clearly in the prompt** (e.g. "run to verify"), and dsh will follow through to completion.

Hands-on exercises

- Reproduce Case A:** create a synthetic JSON file with dirty data (missing values, type errors, duplicates, outliers). Give it to dsh with the same prompt. Compare your results with the case study.
- Reproduce Case B:** create a Python module with 5 planted bugs. Ask dsh to find and fix them, then write tests. How many bugs does it find? How many tests does it write?
- Prompt variation:** run Case A twice. First time, say "clean the data." Second time, say "clean the data, run and verify, and explain your trade-offs." Compare the output quality.
- Timing analysis:** measure how long each step takes (read, write, bash, verify). Which step is the bottleneck? How does this compare to the 90% / <1% split from Chapter 6?
- Failure injection:** give dsh an impossible task (e.g. "fix this code, but don't run the tests"). Does it still try to verify? What happens when the acceptance criteria are unclear?
- Think:** why does dsh "show judgment" (explain trade-offs, design edge-case tests)? Is this a feature of the model, the harness, or the prompt?

FAQ

- Q:** Are these real-world cases or synthetic? Synthetic. The data and code are self-authored to demonstrate capabilities. No real business data is involved.

- **Q: Why 94 seconds for Case B?** That's the total wall-clock time for reading the code, finding 5 bugs, fixing them, writing 49 tests, running them, and summarizing. Most of the time is model thinking (Chapter 6).
- **Q: What if dsh misses a bug?** In Case B, it found all 5. But in general, the quality depends on the prompt. If you say "find all bugs," it will try. If you say "fix the obvious bugs," it might miss subtle ones.
- **Q: Can I use dsh for production data cleaning?** Yes, but review the output. dsh explains its trade-offs (e.g. "median imputation creates a spike"), but you should verify the cleaning logic matches your business rules.
- **Q: What's the "artifact tracking" mentioned in the summary?** Tool returns carry ``locations`` (file paths). dsh uses these to build "artifact file lines" — the artifact chips at the end of the conversation. You can click to open each file.
- **Q: How do I handle failure recovery in production?** Pair dsh with guard timeouts (to prevent infinite loops) and human approval (for sensitive operations). The ``guard/`` and ``interaction/`` packages provide these capabilities.

Appendix A: [Glossary & Command Quick Reference](./appendix-glossary.md)

Chapter 11: Future Outlook — Where dsh and the Agent Ecosystem Are Headed

> Goal: project the future of dsh and its ecosystem from multiple angles — *technology, ecosystem, competition, industry, developer opportunity, risk*. These are projections based on current facts, not assertions — use them to judge whether it's worth investing.

TL;DR (30-second version)

- 1. **Short term (1–3 months)**: official 0.1.0 release, plugin ecosystem boom, tutorial/tool projects grab the first wave of dividends
- 2. **Mid term (1 year)**: dsh becomes a candidate de-facto standard for the "programmable agent base", vertical plugins mature
- 3. **Long term (2–3 years)**: agent engineering layer standardizes; dsh ecosystem and model iteration reinforce each other
- 4. **Biggest risks**: breaking changes, ecosystem fragmentation, official strategy shifts
- 5. **Individual opportunity**: entering now (tutorials/plugins/tools) is cheap early-mover dividend

<details> <summary>Chapter navigation</summary>

- [11.1 Technology projections](#111-technology-projections)
- [11.2 Ecosystem projections](#112-ecosystem-projections)
- [11.3 Competitive landscape](#113-competitive-landscape)
- [11.4 Industry applications](#114-industry-applications)
- [11.5 Developer opportunities (those who enter now)](#115-developer-opportunities-those-who-enter-now)
- [11.6 Risks & uncertainty (honest edition)](#116-risks--uncertainty-honest-edition)
- [11.7 This handbook's own outlook](#117-this-handbooks-own-outlook)
- [11.8 Final advice to readers](#118-final-advice-to-readers)

</details>

11.1 Technology projections

Horizon	Projection	Basis	
Short	0.1.0 stable release drop rc, breaking changes converge	Currently rc.6; official explicitly warns of breaking changes;	

		iteration is fast	
Short	Official TUI & interactive CLI highest community demand, #172 15+ comments	Hottest request in official Discussions	
Mid	Cache/performance optimization becomes the theme finer reasoning effort, cache-hit tooling	High cache-hit rate is already a cost advantage ch. 5; it will be made explicit via tooling	
Mid	Plugin ecosystem standardization marketplace / one-click install / quality tiers	Official dsh-plugin topic exists; plugin growth inevitably creates a marketplace	
Long	Agent runtime architecture solidifies as an "engineering paradigm" an Agent-era analog of the Linux kernel model	"Everything is a plugin" has proven to be a composable architecture	

11.2 Ecosystem projections

- **Plugin count**: day zero → tens to hundreds within a year (referencing similar ecosystem S-curves)
- **Content ecosystem**: tutorials/blogs/awesome lists erupt first (Chinese content is currently a vacuum — this handbook is the first systematic tutorial)
- **Community size**: official Discussions adds dozens of threads daily (we observed 30+ new topics in a single day); Discord activity will grow in tandem
- **Flagship projects**: a few "ecosystem flagships" will emerge (think Prettier/ESLint status in the VSCode ecosystem) — **tool-type** and **tutorial-type** are most likely

11.3 Competitive landscape

Dimension	How dsh's positioning evolves	
vs Claude Code	Short term still "less out-of-the-box"; long term differentiates on customization + open source + cost	

vs OpenCode	Both open source; dsh's official backend + plugin ecosystem is the differentiator OpenCode has no official backend	
vs building your own	dsh is the bridge from "half-finished to finished" — less work than DIY, more control than closed source	
Model side	Each DeepSeek model generation amplifies dsh's ecosystem value model + runtime synergy	

****Key judgment****: dsh's winning move isn't "how strong the model is" (that's DeepSeek's job) — it's whether the ****agent engineering layer can become the standard****. Whoever defines "how plugins are written, how the ecosystem grows" wins that layer.

11.4 Industry applications

- ****Finance/data analysis****: long sessions + high cache-hit = cost-friendly batch analysis, first to land
- ****Enterprise knowledge/docs****: 1M context + document toolchain makes internal knowledge management a natural fit
- ****Education****: batch problem generation/explanation/summarization — cost-sensitive scenarios eat the cache dividend
- ****Vertical agent products****: "industry agent wrappers" will emerge (industry plugins/config bundles built on dsh)

11.5 Developer opportunities (for those who enter now)

Role	Opportunity	Timing	
Tutorial/content authors	Chinese tutorial vacuum this handbook is the first	Now biggest dividend	
Plugin developers	Directions the official repo lacks: TUI, remote access, mobile, cache tooling	Now	
Tool-type projects	Speedup, latency visualization, MCP	Now we validated with the speed-up plugin	

	toolkits	experiment	
Ecosystem infrastructure	Plugin marketplace, template generators, benchmark tooling	Mid term	
Industry solutions	Finance/education industry plugin bundles	Mid term	

11.6 Risks & uncertainty (honest edition)

1. **Breaking changes**: rc-stage API may change frequently — plugins/tutorials need ongoing follow-up (this handbook updates in sync)
2. **Ecosystem fragmentation**: plugin quality varies, standards missing — early stage needs "high-quality showcase projects" (this handbook + the Chapter 4 example are showcases)
3. **Official strategy shifts**: the official team may pull back directions (e.g., built-in TUI) — ecosystem projects need differentiated moats
4. **Model competition**: other models' iterations may weaken DeepSeek's appeal — but dsh's model-agnostic design (OpenAI-compatible endpoints) is a buffer
5. **Heat cooling**: if the dsh ecosystem doesn't take off, early investment may be "wasted" — but tutorials/skills themselves are transferable assets

11.7 This handbook's own outlook

- **Content**: continuously updated as dsh iterates (version-aligned with rc/stable releases)
- **Bilingual**: English edition catches up with Chinese
- **Cases**: more real industry cases (compliance-first)
- **Ecosystem**: link with awesome lists and official Discussions to become one of the ecosystem's content hubs

11.8 Final advice to readers

- **Want to try it**: install and play now (ch. 2) — cost is near zero
- **Want to invest**: pick a direction that is "missing from the official repo + matches your strengths" (TUI / industry plugins / tutorials)
- **Want to wait**: wait for the 0.1.0 stable, but you'll miss the early-ecosystem dividend — **early dividend $\approx$ patience of the first movers**
> Predictions can be wrong, but "entering now costs the least" is probably not.

Chapter 12: Known Limitations & Boundaries — the Honest Edition

> *Goal: list dsh's **known limitations** and usage boundaries without hype. When this handbook was written, dsh was still at `0.1.0-rc.6` (pre-release). These limitations are "facts of this moment", not "permanent fate" — they will change as versions iterate.*

>

> *All "known issues" come from hands-on testing + public feedback in the official repo discussions (deepseek-ai/deepseek-harness, as of 2026-08-13).*

TL;DR (30-second version)

1. **The rc stage is the biggest uncertainty**: API/config can break at any time — pin your version line before investing
2. **Plugin ecosystem is early**: 60+ official packages but limited coverage; few third-party plugins; interactive forms (TUI) missing
3. **Cross-platform gaps**: real path/encoding bugs on Windows (UTF-16, 0x00)
4. **Performance boundaries unverified**: high concurrency, large codebases lack public stress data
5. **Official strategy risk**: open-source commitment vs commercialization is the black swan you must consider when choosing

12.1 rc-stage instability

Breaking changes happen anytime

dsh has gone through at least one **dependency breakage** from rc.1 to rc.6: inconsistent `@deepseek-ai/\*` version lines caused plugins to fail installing (404/ERR_MODULE_NOT_FOUND). rc.6 has converged a lot, but **nothing guarantees rc.7 or the stable release won't break again**.

Real impact on you:

- Tutorials/plugins may need rewriting every couple of weeks
- Pinning the version line (`^0.1.0-rc.6`) is the baseline; but whether an rc-line upgrade is "backward compatible" or "start over" depends on the official team's mood

> *[!WARNING]*

> *Evaluate carefully before production use. At writing time the official team is iterating fast — **what works today may be deprecated tomorrow**.*

Known rc pitfalls (reproduced locally / in the community)

Pitfall	Symptom	Status	
link-dev dependency breakage	@deepseek-ai/ ERRMODULENOTFOUN D when linking locally, but works via npm install	Root-caused flat fallback dir mechanism; pin the dependency line to avoid	
Windows path 0x00	workspace creation fails with Chinese/special-char paths ENOENT	Community-reported #151/#396	
fs-sandbox race	post-check path check races with workspace- write	Community-reported #159	
TUI missing	no interactive terminal UI in official examples community builds exist, not merged	Proposed #392	

12.2 Early plugin ecosystem

****Current state****: 60+ official packages, but the ecosystem is overall at "day zero".

- ****Few third-party plugins****: plugin threads in Discussions are exploding, but quality varies — most are first-plugin practice pieces
- ****Missing interactive forms****: headless/JSON-RPC/Web are complete, but ****no official example of terminal-native UI (TUI)**** (the community deepseek-harness-tui proved it viable; not merged into official)
- ****Docs lag the code****: official docs are architecture-focused; onboarding paths and pitfall records depend on community output (which is exactly why this handbook exists)
- ****Version-line chaos****: multiple version lines coexist on npm; installing the wrong line = 404

****Judgment****: entering the ecosystem now is early-mover dividend (less competition, official willingness to support), but ****don't expect "just follow the official docs and it runs"***** — hitting pitfalls is the norm.

12.3 Cross-platform gaps

dsh is mature on macOS/Linux, but Windows support has a clear "second-class citizen" feel:

- **Path handling**: real bugs with Chinese/special-character paths (0x00 truncation, ENOENT); multiple community reports
- **Encoding**: UTF-16 handling has defects; occasional garbled terminal output
- **CLI/TUI debate unsettled**: there's an ongoing dispute in official Discussions about "CLI is the way vs TUI is the interactive future" (#386); investing in a terminal-UI form before the direction settles carries risk

12.4 Unverified performance boundaries

The handbook's benchmark covers **single agent, normal task scale**. These scenarios lack public data:

- **High concurrency**: parallel agents, large-scale task distribution
- **Large codebases**: context management and tool-call latency at million-line scale
- **Long-running stability**: hours-long sessions, memory-leak risk
- **Cache-hit boundaries**: the measured ~97% hit rate depends on "repetitive conversation patterns" — brand-new/cold-start tasks drop noticeably

12.5 Gaps vs mature agents

Compared with Claude Code, Codex, and other mature products, dsh's gaps are **structural** (not fixed by adding a few features):

Dimension	dsh rc.6	Mature agents	
Ecosystem stage	day zero, few third-party assets	mature, complete	
Out of the box	requires understanding profiles/plugin system	install & run	
Stability	rc stage, breaking-change risk	stable releases, long-term compatibility promises	
Docs	official is architecture-focused; onboarding via community	complete official docs + rich tutorials	
IDE/toolchain integration	early	deep	

dsh's positioning: not an "out-of-the-box car", but a "programmable LEGO base".
Choosing dsh = choosing **freedom**, at the cost of **assembling it yourself**.

12.6 Official strategy risk

Black swans a selector must consider:

- **Open-source commitment**: MIT with no noise so far, but "official-grade plugin system" means the ecosystem is deeply bound to the official cadence
- **Model-binding tendency**: default experience favors DeepSeek models (any model via the llm plugin, but "official optimization" is DeepSeek-first)
- **Commercialization shift**: once dsh user numbers rise, the official team may launch managed services/paid tiers — whether the open-source edition keeps parity is unknown
- **Direction drift**: rc-stage features are added/removed frequently (the CLI/TUI debate is a signal); betting on the wrong direction = sunk cost

12.7 This handbook's own limitations

Finally, honestly about this handbook itself:

- **Written against rc.6**: every command, config, and data point verified on rc.6 — **newer versions may invalidate everything**
- **Limited benchmark sample**: single machine, single model, limited task set — not an authoritative evaluation
- **Chinese-first**: the English edition is a translation/condensation and may lag the Chinese content
- **Incomplete coverage**: dsh has 60+ official packages; this handbook deep-dives the core path; long-tail plugins aren't covered one by one
> **Usage advice**: treat this handbook as an "rc.6 snapshot + methodology", not an "eternal bible". When versions update, run through the [roadmap learning path](./roadmap.md) acceptance checks first, then decide whether to update your dependency line.

Info as of 2026-08-13 (dsh 0.1.0-rc.6). For newer versions, issues/PRs are welcome to correct.

[English](./13-security.en.md) | [中文](./13-security.md) · [← Back](../README.md)

Chapter 13: Security & the Sandbox Model

> ***Goal of this chapter:** Understand dsh's security skeleton — what the sandbox governs, how permissions are tiered, how approvals gate dangerous actions — and the real boundary cases the community has found. **This is the chapter that takes you from "it runs" to "safe enough for production."***

TL;DR (30-second version)

1. **Three layers of defense**: sandbox (where it executes) → permission presets (what it can do) → approval (whether dangerous actions need a human nod) — execution is *not* naked by default
2. **The sandbox only governs file effects**: ``read-only`` / ``workspace-write`` / ``danger-full-access``; network and process visibility are outside its vocabulary (so cases like ``taskkill`` killing the host are "outside the vocabulary", not a sandbox breach)
3. **Permissions = two knobs bundled into presets**: ``sandbox/mode`` + ``approval/policy``; two default presets: ``workspace-write`` (workspace writes + ask) and ``danger-full-access`` (no sandbox + never)
4. **Approval is fail-closed**: only ``allowed-once`` grants; ask with no responder = rejected (``unavailable``), never rejects deterministically — the default headless posture
5. **Community audits found real boundaries**: #159 fs-sandbox race, #584 scrubbedParentEnv collateral damage, #381 iframe clickjacking, #817 audit report (vm escape + unauthenticated local RPC)
6. **Enterprise baseline in one line**: default ``workspace-write`` + ask, keep loopback, install plugins only from trusted sources, keep sensitive directories out of the workspace

<details> <summary>Chapter navigation</summary>

- [13.1 Security Design Philosophy: The Three Layers](#131-security-design-philosophy-the-three-layers)
- [13.2 Sandbox Mechanics: FS Boundaries and Tool Sandboxing](#132-sandbox-mechanics-fs-boundaries-and-tool-sandboxing)
- [13.3 Permission Model: Tool Permission Tiers](#133-permission-model-tool-permission-tiers)
- [13.4 Approval Flow: From Request to Grant](#134-approval-flow-from-request-to-grant)
- [13.5 Plugin Security Audit Checklist](#135-plugin-security-audit-checklist)

- [13.6 Known Security Cases from the Community](#136-known-security-cases-from-the-community)
 - [13.7 Enterprise Security Baseline](#137-enterprise-security-baseline)
- </details>

13.1 Security Design Philosophy: The Three Layers

dsh's security skeleton is three layers, each with one job (based on the official ``docs/subsystems/*`` architecture docs, verified: ``sandbox.md`` / ``permission-presets.md`` / ``approval.md``; cross-checked against the hands-on measurements in Section 8.5):

Layer	What it governs	Official packages	One-liner	
Sandbox	Where code runs, which files it can touch	sandbox/sandbox, sandbox/sandbox-local, sandbox/sandbox-policy	Isolates the file effects of command execution	
Permissions	What the agent may do	interaction/permission-presets	Bundles sandbox + approval into named presets	
Approval	Whether dangerous actions need a human confirm	interaction/user-approval	The last gate before least privilege	

****Least privilege**** is the shared design intent across all three layers: the default preset only grants "workspace writes + ask before critical operations"; any escalation (switching to ``danger-full-access``, changing the approval policy) happens explicitly and is written to the session log — the ``permission/preset`` event is a ****log-only record of user intent****, kept out of the model transcript, so you can always look back at "who raised privileges, and when."

****Threat model in one line****: the three layers defend against "untrusted input driving the agent across its boundaries" — prompt injection, malicious plugins, induced high-privilege operations (#381 iframe hijacking is a direct shot at this chain). They do *not* promise to defend against "trusted code going rogue" (a plugin *is* executable code, see 13.5), nor against network/process-level attacks.

> *Core takeaway: ****dsh's tool execution has isolation and approval layers by default — it is not naked execution.**** But it is not a "security operating system" either — the sandbox doesn't block network or processes, and approval can be self-answered by the model (13.6).*

13.2 Sandbox Mechanics: FS Boundaries and Tool Sandboxing

The official sandbox doc is explicit: **the sandbox only governs file effects (filesystem effects)** — network and process visibility are not part of its vocabulary.

The three modes (`SandboxMode``):

Mode	Allowed file effects	Notes	
read-only	Read-only + required sinks e.g. /dev/null on POSIX	Windows ACL backend has no explicit writable root → reports partial	
workspace-write	Workspace root + backend-promised temp areas writable	Default preset; root comes from the session's immutable cwd	
danger-full-access	Unrestricted	Spawns the original argv directly, bypassing <code>ctx.sandbox</code>	

Execution details:

- Policy is per-call:** `SandboxExecutionPolicy`` (mode + workspaceRoot + sessionId) is resolved at each capability call rather than being pinned to a provider — the same provider can serve a `read-only`` bash and a subagent that needs a writable state directory
- Workspace root canonicalization:** first canonicalize per filesystem semantics (`symlink/..`` resolved to real directories), then apply lexical normalization — prevents symlink-based boundary escapes
- Backend matrix:** Linux bwrap/Landlock, macOS Seatbelt, Windows ACL restricted-token; `enforcement: full | partial`` — old Landlock ABIs and the Windows ACL Everyone/hard-link boundaries are the current `partial`` cases (**when the promise is watered down, consumers must treat it as `partial``**)
- FS tools share the boundary:** file tools like `write`/`edit`` are held to the same `workspace-write`` constraint — it's not "file tools pass through, only bash is sandboxed" (#149's recursive workspace deletion happened under `workspace-write``)
- Writes restricted, reads not:** the current permission model constrains **writes**; content outside the workspace can still be read (#492 raised this when proposing an evaluation isolation mode)

****Tool sandboxing****: ``bash`/`pwsh`` go through ``bash-sandbox` / `pwsh-sandbox`` (consuming ``ctx.sandbox``) and spawn child processes inside the sandbox. On Windows, the pwsh sandbox has ACL constraints (we hit temp-permission issues in Chapter 8's hands-on testing; #758 also reports that cleaning up the Windows sandbox temp directory can permanently crash it).

****Related family cases****: #159's race is a TOCTOU (the path is swapped between check and use); #278 is a different widening trick — when ``/tmp`` is the workspace, a restricted child process can use rebinding to widen the ``workspace-write`` grant. Shared lesson: ****choose a trusted location for the boundary root, and never rely on a single path check alone.****

13.3 Permission Model: Tool Permission Tiers

The official permission presets bundle ****two knobs**** — ``sandbox/mode`` and ``approval/policy`` — into named presets, surfaced to the client as a single "Permissions" selector. Two default presets:

Preset	sandbox/mode	approval/policy	Semantics	
workspace-write	workspace-write	ask	Writable inside workspace, asks before operations default	
danger-full-access	danger-full-access	never	No sandbox, no asking	
custom derived state	any combination	any	The "non-preset" state after manually turning the knobs; cannot be a switch target	

****Tool permission tiers**** (ordered by blast radius; "suggested preset" per tool is derived from the official docs, not tested tool-by-tool — ****inferred, pending verification****):

Tier	Representative tools	Suggested preset	Risk	Notes	
Read-only probing	read/grep/glob	read-only compatible	Low	Read-only, nothing written; reads not limited to the workspace	

Workspace writes	write/edit/streplaceditor	workspace-write	Medium	Write boundary = workspace root + temp areas	
Command execution	bash/pwsh	workspace-write sandboxed	High	The sandbox doesn't block network or processes	
Full access	any switch to danger-full-access	danger-full-access	Extreme	All gates off	

> The official docs are explicit that the sandbox **only governs files** — so the high risk of the "command execution" tier can't be sandboxed away. It must be backstopped by approval and trust boundaries (13.5 / 13.7).

****Real boundary cases**** (details in 13.6): recursive deletion of the entire workspace with zero confirmation under ``workspace-write`` (#149); ``danger-full-access`` accidentally deleting an entire home directory (#461, real incident); on Windows, the minimal preset allows out-of-workspace writes with no approval (#523); an agent inside the sandbox can ``taskkill`` the host (#466 — process visibility is outside the sandbox vocabulary).

13.4 Approval Flow: From Request to Grant

The official approval doc: approval answers one question — "can this specific action happen?"

****Fail-closed result set**** (``ApprovalOutcome``):

Outcome	Meaning	Effect on the caller	
allowed-once	One-time grant only for the action asked about	✅ proceed	
rejected	Explicit denial	❌ abort	
cancelled	Request withdrawn AbortSignal	❌ abort	
unavailable	No responder / responder errored	❌ abort fail-closed by default	

****Policy**** (``ApprovalPolicy``): ``ask`` (default — handed to the responder chain; empty chain → ``unavailable``); ``never`` (deterministically rejects everything, dispatches no responders — the strict headless/CI posture).

****Flow****: tool call → `approval/request` waterfall (plugins can intercept/rewrite) → responder (a human in the Web UI, or a one-shot machine decision via the ACP bridge) → every request writes a paired audit pair `approval/asked` + `approval/decided` (same `ApprovalRequestId`).

****Request content**** (`ApprovalRequest`): agent (a responder only answers for the agents it owns) + toolName + callId (linked to the already-streamed tool call, so parameters aren't re-rendered) + reason (why the question is being asked) + signal (AbortSignal for withdrawal).

****Three known pitfalls****:

1. ****Approval loop****: #250 reproduced "model self-approving danger-full-access" for real (the Web approval channel can be driven by the model itself) — ****approval is not a security boundary****, only a human-in-the-loop gate
2. ****Concurrent misrouting****: #453 — clicking Allow during concurrent sandbox approvals aborts a *different* call (UI response and callId mismatched)
3. ****Reconnect drops pending approvals****: #646 — silent loss of pending approvals on reconnect (resync clears them + replay race); headless approval behavior is undefined (#291)

13.5 Plugin Security Audit Checklist

Plugins are dsh's capability boundary ("everything is a plugin", see Chapter 4) — ****installing a plugin = installing executable code.**** The checklist distilled from community audits (#817 and family):

#	Check	Community basis	Disposition	
1	Is the source trustworthy? dsh plugin add has no signature/source validation	#587	Install only from official/well-known sources; read the source before installing	
2	Minimal boot-time permissions? Plugins get write access to the whole config tree at boot	#587	Beware plugins that "auto-edit your config after install"	
3	Not treating the vm as a security boundary? Workflow/dynami	#243 #451 #774 #778	Treat the vm as an isolation engine, not a security layer	

	c-plugin vms can escape			
4	Child-process env scrubbed correctly? scrubbedParentEnv substring-matching hits legitimate variables	#584	Check whether variables containing the KEY substring e.g. KEYBOARD/MON KEY get wrongly scrubbed mechanism inferred from the thread title	
5	Secret redaction not fail-open? settings role'secret' has redaction gaps	#226	Don't put secrets into prompts; verify the redaction logic	
6	Security hooks not silently disabled? timeout: 0 fails open	#460 #583	Set positive hook timeouts; loading failures must surface as errors	
7	Audit trail? approval/asked/decided pairs are queryable	official approval.md	Critical operations traceable per session	
8	No plaintext session leakage? .tmp plaintext session residue after crashes	#674	Disable crash dumps in sensitive environments / clean up periodically	

13.6 Known Security Cases from the Community

All 4 threads below verified to exist via `gh api graphql` (deepseek-ai/deepseek-harness Discussions, 2026-08-14):

Thread	One-liner	Type	Impact	
#159 https://github.com/deepseek-ai/deepseek-harness/	fs-sandbox post-check pathname race can bypass the workspace-write file	Sandbox race	TOCTOU: the path is swapped between check and use, enabling out-of-bounds	

discussions/159	boundary		writes	
#584 https://github.com/deepseek-ai/deepseek-harness/discussions/584	scrubbedParentEnv substring-matching wrongly scrubs legitimate env vars like KEYBOARD/MONKEY cherry-pickable fix included	Child-process env	Legitimate env vars get scrubbed → behavioral anomalies	
#381 https://github.com/deepseek-ai/deepseek-harness/discussions/381	Default localhost web UI can be clickjacked via cross-site iframe, inducing a Full access grant and driving high-privilege operations	Frontend clickjacking	Malicious webpages trick users into clicking out high privileges in the dsh UI	
#817 https://github.com/deepseek-ai/deepseek-harness/discussions/817	Security audit report: sandbox/approval boundary bypasses + unauthenticated local RPC PoC available privately	Systematic audit	Covers vm escape family #243 #778, unauthenticated local /api family #451, CVSS 8.8, approval loop #250	

****High-value family supplements**** (thread numbers from `docs/research/discussion-mining.md` §2.6 in this repo; that report says all 780 threads were programmatically verified):

Thread	One-liner	Lesson	
#149	Whole-workspace recursive deletion with zero confirmation under workspace-write	Workspace contents also need protection from "self-destruction"	
#250	Web approval loop: model self-approving danger-full-access	Approval is not a security boundary	
#461	Full Access mode deleted an entire home directory real incident	High-privilege preset risk, proven in the wild	

#587	Third-party plugins get whole-config-tree write access at boot	Plugin trust = supply-chain security	
#466	An agent inside the sandbox can taskkill the host	Process visibility is outside the sandbox vocabulary	
#674	.tmp plaintext session residue after crashes is never cleaned	Privacy risk	

> Disclaimer: these are community reports from the rc.6 era (within 48h of the 2026-08-13 release); some may already be fixed in newer rcs — keep the version context in mind when citing.

13.7 Enterprise Security Baseline

Decision guidance for "going to production / rolling out to a team" (aligned with Chapter 6's style):

Dimension	Baseline	Basis	
Network exposure	Keep loopback official rejection of --host 0.0.0.0 is intentional; don't bypass until remote auth matures	#76 #130 #397	
Permission preset	Default workspace-write + ask; danger-full-access only on isolated machines / one-shot tasks	official preset table + #461	
Approval policy	Human responder chain required; headless uses never deterministic rejection with a review gate in front	approval.md + #291	
Plugin governance	Whitelisted sources + run the 13.5 checklist before installing; ban bare dsh plugin add github:	#587 #656	
Workspace boundary	Sensitive directories home/keys stay out of	#149 #461	

	the workspace; workspace contents go under version control regularly		
Audit	Session logs retained approval pairs traceable; guard against .tmp plaintext residue	approval.md + #674	
Upgrade discipline	Re-run the 13.5 checklist on every rc upgrade sandbox backends / preset table can change	Chapter 12 rc-integrity stance	

****Decision memo****: local personal dev → default preset is enough; team-shared → `workspace-write` + ask + plugin whitelist; CI/unattended → headless + `never` + code review gate in front; don't run `danger-full-access` long-term anywhere just for convenience.

****Scenario → recommended combination**** (decision guidance, derived from the baseline table above):

Scenario	Recommended combination	Why	
Local personal dev	Default workspace-write + ask	Works out of the box; approval isn't annoying	
Team-shared host	workspace-write + ask + plugin whitelist	Multi-user; plugin supply-chain risk is amplified	
CI / unattended	headless + never + review gate in front	Deterministic rejection; nothing hangs waiting on a human	
Isolated machine / one-shot task	danger-full-access the machine itself is isolated	Maximum authorization, but blast radius capped at one box	

Hands-on exercises

1. ****Understand****: what file effects does each of the three sandbox modes allow? Why is "network and processes are outside the sandbox vocabulary"?

- > Check yourself: Section 13.2's three-mode table and the "backend matrix" paragraph
 - 2. **Understand**: why is ``never`` under headless the "strictest" rather than the "loosest"?
- > Check yourself: Section 13.4's "Policy" paragraph (deterministic rejected)
 - 3. **Hands-on**: open ``dsh web``, look at the two presets in the Permissions selector; switch to ``danger-full-access`` and back, and watch the ``permission/preset`` event in the session log
- > Check yourself: Section 13.3's preset table
 - 4. **Hands-on**: before installing a community plugin, walk it through the 13.5 checklist item by item (source/source code/permissions/audit) and write the conclusion into your notes
- > Check yourself: Section 13.5's checklist table
 - 5. **Think**: why does #250's "model self-approval" prove that approval is not a security boundary? If approval can't stop it, what does an enterprise fall back on?
- > Check yourself: Section 13.6 + the 13.7 baseline table

FAQ

- **Q**: Can the sandbox stop malicious code? **No** — don't use it as a security boundary. The sandbox only governs file effects; network (e.g. SSRF) and processes (e.g. #466's ``taskkill``) are outside its vocabulary. Defending against malicious code relies on plugin trust (13.5) and network/process-level isolation (containers/VMs/remote sandboxes — the official docs are explicit these are sibling implementations of the capability seam, not ``ctx.sandbox`` providers).
- **Q**: Why is ``never`` under headless actually safer? ``never`` isn't the lax "never ask" mode — it's **deterministic rejection**: any approval request is directly ``rejected``, with no responders dispatched — unattended runs never "hang waiting for someone to click Allow." Actions that need to happen rely on an upfront review gate or explicit configuration, not runtime questions.
- **Q**: So is ``workspace-write`` safe? Relatively safe inside the boundary — but #149 shows "recursively deleting the whole workspace" gets zero confirmation under ``workspace-write``. The sandbox stops *boundary crossing*, not *self-destruction inside the boundary*. Put important code under version control.
- **Q**: Chapter 12 already mentioned #159 — why cover it again here? Chapter 12 was a "known-issues quick table" (one-liner + status); this chapter expands the mechanism (how the post-check race happens), the family (#278 /tmp

rebind), and the mitigations (don't trust a single path check, keep critical directories out of the workspace).

- **Q:** Can these security thread numbers be trusted? Surely someone could have made them up? **A:** The 4 core threads (#159 #584 #381 #817) were each verified to exist via `gh api graphql` (links in the 13.6 table); the family supplements come from docs/research/discussion-mining.md` in this repo (780 threads fetched with full pagination and programmatically verified).`

Next chapter: [\[Chapter 14: Caching & Cost Engineering\]\(./14-cost.en.md\)](#)

[\[English\]\(./14-cost.en.md\)](#) | [\[中文\]\(./14-cost.md\)](#) · [\[← Back\]\(../README.md\)](#)

Chapter 14: Caching & Cost Engineering

> ***Goal of this chapter:** Turn "it's cheap" from folklore into engineering — the conversation-cache mechanism, measured hit rates and how to improve them, the cost model, reasoning-tier interplay, real-task budgeting, and how to *see* where every token goes.*

TL;DR (30-second version)

1. **Cache is the #1 cost variable**: DeepSeek's context cache bills repeated input at a **98% discount (Flash) / 99%+ discount (Pro)** — agent workloads are input-heavy, so the hit rate directly determines the cost scale
2. **Measured hit rate in this handbook: 97%**: session continuity + stable tool schemas + the nature of multi-step agent workloads (Section 5.1.1); community long-run measurements reach 99.7%
([#560](https://github.com/deepseek-ai/deepseek-harness/discussions/560))
3. **Three hit-rate principles**: keep the session alive (don't keep creating new ones), keep the prefix stable (don't keep changing config), interrupt less (write clear acceptance criteria once)
4. **Two independent levers**: reasoning tier controls *how much thinking*, cache controls *the input unit price* — they stack, so `low` + high hit rate = the cheapest combination`
5. **Cost must be visible**: the session stats line shows cache hit %; per-turn token counts are awaiting official implementation
([#735](https://github.com/deepseek-ai/deepseek-harness/discussions/735)) — until then, use the stats line + a community cost-tracker plugin

<details> <summary>Chapter navigation</summary>

- [14.1 How Conversation Caching Works](#141-how-conversation-caching-works)
- [14.2 Measured Hit Rate: Where 97% Comes From](#142-measured-hit-rate-where-97-comes-from)
- [14.3 Optimizing the Hit Rate in Practice](#143-optimizing-the-hit-rate-in-practice)
- [14.4 Cost Model: How the Hit Rate Drives Your Bill](#144-cost-model-how-the-hit-rate-drives-your-bill)
- [14.5 Reasoning Tiers and Caching: the Cost Matrix](#145-reasoning-tiers-and-caching-the-cost-matrix)
- [14.6 Budget in Practice: Cost Estimate for Case A](#146-budget-in-practice-cost-estimate-for-case-a)

- [14.7 Measurement: Making Token Usage Visible](#147-measurement-making-token-usage-visible)
 - [14.8 Cost Pitfall Checklist](#148-cost-pitfall-checklist)
 - [14.9 Decision Recommendations at a Glance](#149-decision-recommendations-at-a-glance)
- </details>

14.1 How Conversation Caching Works

DeepSeek's API **context cache** (prompt caching): repeated/similar input tokens are billed at the **cache price** instead of full price. dsh's input for every model request = system prompt + skill catalog + conversation history + tool schemas (Section 8.3) — **the stable prefix portion is the natural cache candidate**, which is why agent workloads have far higher cache potential than ordinary chat.

Billing rules (Section 5.1.1 pricing table):

Charge	Miss price	Cache-hit price	Discount	
Input Flash	\$0.14/M	\$0.0028/M	98%	
Input Pro	\$0.435/M	\$0.003625/M	99%+	

Input structure per request (Section 8.3) — determines which tokens can hit:

system prompt (stable) + skill catalog (stable) + conversation history (incremental) + tool schemas (stable) + this turn's messages (incremental) cache-hit zone (repeated) miss zone (new increments)

The bigger the hit zone and the smaller the increments → the higher the hit rate. Once you understand this structure, the three optimization levers in 14.3 are all really "make the hit zone bigger."

Cache isn't just cheaper, it's faster: cold start with context injection takes ~110s on the first turn; a hot cache takes ~1s (measured in Chapter 1) — first-token time drops by an order of magnitude on hits.

> Detail note: hits depend on "repeated/similar prefix"; the exact matching rule is decided API-side (inferred, pending verification). In practice, "keep the prefix stable" is enough to reliably capture the discount.

14.2 Measured Hit Rate: Where 97% Comes From

This handbook measured a cache hit rate of up to **97%** in real dsh session stats (Section 5.1.1). Why dsh's hit rate is unusually high:

Reason	Explanation	
--------	-------------	--

Session continuity	Multiple turns in one session repeat the system prompt / skill catalog / history → hits	
Stable tool schemas	Every step carries the same tool definitions → hits	
Agent workload nature	Multi-step tool chains repeatedly carry the same context → naturally high hit rate	

****Boundary (honest edition)**:** 97% depends on "repetitive conversation patterns" — ****brand-new tasks / cold starts drop noticeably**** (Section 12.5). Community long-run measurements reach ****99.7%**** ([#560](https://github.com/deepseek-ai/deepseek-harness/discussions/560)), confirming "longer sessions → higher hit rates."

****How it was measured**:** the "cache hit %" in the session stats line at the bottom of the Web UI (Section 6.6); the comparison baseline is 5.1.1's "50-step tool chain task, 2.4M input tokens" — a typical sample of the input scale.

Metric	Where to look	Behavior on hits	
Cache hit %	Session stats line bottom of Web UI	Stays high 97%+ in long sessions	
First token	Per-turn end line	Drops by an order of magnitude from 10s+ to under 1s FAQ Q4 in Chapter 6	
Total tokens	Session overview	Small increments = the hit zone is growing 14.1 structure	

14.3 Optimizing the Hit Rate in Practice

Three levers, ordered by cost-effectiveness:

Lever	What to do	Why it works	
Session continuity	Keep long tasks in one session; avoid creating new sessions frequently; batch tasks in the same session/prefix	Repeated history messages → hits	
Stable prompts	Keep the system prompt / skill catalog prefix stable; don't	Identical prefix → hits	

	churn config; keep tool schemas unchanged		
Fewer interruptions	Write clear acceptance criteria once Section 5.7 / the lesson in Chapter 10's cases instead of "reopening sessions and re-explaining context"	Avoids rebuilding the prefix → full-price restart	

****Grounding check****: when the hit rate drops below expectations, the first thing to suspect is "did the prefix change?" — changing config, switching tiers, or touching the skill catalog all invalidate the cache (Section 6.6 monitoring line).

14.4 Cost Model: How the Hit Rate Drives Your Bill

****Cost model in one line**** (Section 6.6): ``total cost ≈ output tokens × output price + missed input × miss price + hit input × hit price`` — agent workloads are input-heavy, so ****the cache hit rate is the #1 cost variable****.

Using Chapter 5's Flash input prices, here's a worked example (****10k input tokens****, hit-rate comparison). Arithmetic sample (90% hits): ``9,000 × $0.0028/M + 1,000 × $0.14/M ≈ $0.000165`` — split hits from misses first, then multiply each by its own price:

Hit rate	Hit tokens	Miss tokens	Input cost Flash	Relative cost	
90%	9,000	1,000	≈ \$0.000165	1×	
50%	5,000	5,000	≈ \$0.000714	≈ 4.3×	
10%	1,000	9,000	≈ \$0.001263	≈ 7.6×	

Scaled to 1M input tokens: 90% hits ≈ ****\$0.0165****, 10% hits ≈ ****\$0.126****, full price \$0.14 — ****nearly an order of magnitude apart****.

Scaled again to this handbook's measured scale (5.1.1: 50-step tool chain, 2.4M input tokens): 97% hits ≈ ****\$0.017**** vs full price ≈ ****\$0.34**** — roughly ****20×****. That's the numerical source of 5.1.1's "cost far below the naive miss-price estimate."

> *Note: this table only covers the input side — Chapter 5's pricing table only records input prices; for the output side, check the official pricing page. In input-heavy agent scenarios, the input side *is* the dominant variable.*

14.5 Reasoning Tiers and Caching: the Cost Matrix

Reasoning tiers control "how much thinking," and cache is an **orthogonal lever**: tiers affect total token count, cache affects the input unit price, and they multiply. Fewer thinking tokens → lower total tokens (Section 6.6), and the context-repeated portion enjoys the cache discount too (inferred, pending verification).

Tier	Thinking tokens	Relative cost	Best for Section 6.2	Cost strategy	
low	Fewest	Lowest	Simple/ deterministic turns: file ops, batch jobs, cheap steps in a tool chain	Default tier for batch jobs	
high	Medium	Medium	Everyday agent tasks default	Default quality/cost balance	
max	Most	Highest	Complex reasoning, long-chain planning, debugging	Fine for one-off complex tasks; don't run it in batch loops	

How the levers combine:

- **Long tool-chain tasks (20–50 steps) benefit the most**: the cumulative effect of downgrading thinking per step is significant (FAQ Q2 in Chapter 6) — both levers (tier × hit rate) act on every step
- Simple turns at `low` show almost no quality difference; complex reasoning at `low` may skip critical steps (FAQ Q3 in Chapter 6)
- **Stacked example** (2.4M-input-scale task, Flash prices): `low` + 97% hits ≈ under \$0.017 (fewer thinking tokens → actually less; inferred, pending verification); `max` + repeated cold starts ≈ over \$0.34 — **the same task is 20×+ apart with both levers fully on vs fully off**
- Note: `low/high/max` are the tier names of the gateway this handbook measured on (pi-ai/opencode-go); the default deepseek-official adapter uses `off/high/max` (Section 6.2 note)

14.6 Budget in Practice: Cost Estimate for Case A

Budget using Chapter 10's **Case A (data-quality analysis, 186s)** as the sample: tool chain `read → write(clean.py) → bash → write(visualize.py) → bash → read →

summarize`, ~6–8 steps. Scaled from the 5.1.1 magnitude (50 steps $\approx$ 2.4M input tokens), Case A's input is roughly **500k tokens** (inferred, pending verification):

Scenario	Hit rate	Input cost Flash, estimated	
Session continuity + stable prefix recommended	97%	$\approx$ \$0.0035	
New session mid-way / prefix changed	10%	$\approx$ \$0.063	
Extreme: full price every time	0%	\$0.07	

Step-level breakdown (why session continuity saves money) — only the first turn is a full-price cold start; every later turn misses only its small increment of "this turn's messages + tool results" (inferred, pending verification):

Step	Content	Input tokens inferred, pending verification	Hit status	
1	readsalesdata.json	~50k	Cold start miss	
2–6	write/bash/ write/bash/read/ summarize	~450k	Prefix hit zone, tiny increments	

Conclusion: a real 3-minute task costs under **1 cent** on the input side; the same work, done differently (session continuity vs repeated cold starts), differs by **~18 $\times$** . The output side (two scripts + summary) is orders of magnitude smaller than the input, so total cost stays in the single-digit cents (inferred, pending verification). This is the per-task source of "dsh with V4-Flash costs about 1/10 to 1/30 of Claude" (measured in Chapter 1).

14.7 Measurement: Making Token Usage Visible

What exists today:

- Session stats line at the bottom of the Web UI: `N turns · M steps | LLM Xs · Tool calls Ys | Avg first token ...` (Section 6.3)
- Per-turn end line: `10:14 · took 9m34s · first token 1.7s · 79 tok/s` ([#735](https://github.com/deepseek-ai/deepseek-harness/discussions/735) original-thread screenshot)
- Session overview shows total-token stats; the stats line shows "cache hit %" (Section 6.6)

****The gap**:** ****per-turn token counts****. Official discussion

[#735](https://github.com/deepseek-ai/deepseek-harness/discussions/735) (filed 2026-08-14, title "【友好显示】希望在每轮对话中加入本轮对话 token 消耗量" / "Friendly display: add per-turn token usage to each turn", 2 comments) is exactly this request — confirmed to exist in the thread; not yet implemented officially.

****Community suggestion (already posted to #735's comment thread)**:** the per-turn display should split into ****two numbers**** — ① total tokens this turn (quick glance at spend); ② hit/miss tokens this turn (see the optimization headroom). Looking at total tokens alone misleads: the same 10k tokens costs 7.6× more at a 90% vs 10% hit rate (see 14.4).

****Interim tooling**:** the community `dsh-plugin-cost-tracker` (Sections 9.5/9.6; real-time token-cost tracking, plugin/MCP form) — implemented through Chapter 4's plugin extension points, tallying `usage` on request/session events (inferred, pending verification); or reconcile manually from session logs / the API usage field. Chapter 11 predicts "cache-hit-rate tooling" is the mid-term theme (11.2) — once #735 lands, the cost-transparency loop closes.

14.8 Cost Pitfall Checklist

#	Pitfall	Symptom	Fix	
1	Misreading cache hits as "the model got faster"	A/B comparison skewed by 1s vs 110s Chapter 6 pitfall #6	A/B tests must use a fresh prompt	
2	Creating new sessions constantly burns money	Every task cold-starts; input billed at full price	Keep long/batch tasks in one session	
3	Config changes silently kill the hit rate	Hit % suddenly drops; cost climbs	Monitor the hit % in the stats line; keep the prefix stable	
4	Judging cost by total tokens alone	10k tokens at 90% vs 10% differs by 7.6×	Split hit/miss and look at both #735 suggestion	
5	Budget missing the output side	Input-only math doesn't add up	Add output prices from the official pricing page Chapter 5 only	

			lists input prices	
6	Tier roulette	max on batch jobs doubles the cost	Use low for batch, max only for complex tasks	

14.9 Decision Recommendations at a Glance

Scenario	Recommendation	
Long tool-chain tasks 20–50 steps	high + session continuity + stable prefix; 2.4M-input scale ≈ \$0.017 97% hits	
Batch / simple turns	low + batched in the same session biggest hit bonus	
Complex reasoning / debugging	max quality over cost; cold start is acceptable for one-off tasks	
Cost-sensitive batch processing	low + headless + serial in one session	
Performance/cost comparison tests	Fresh prompt + fixed tier + median of multiple runs Sections 6.4/6.6	
Hit rate dropping abnormally	Check whether the prefix changed first config / skill catalog / model tier	

Hands-on exercises

1. **Understand**: why is "cache hit rate the #1 cost variable"? Use the 14.4 10k-token table to explain how many times more 90% costs vs 10% hits
- > Check yourself: the 14.4 comparison table (7.6×) and the 6.6 cost-model one-liner
2. **Hands-on**: run a 3-step task and watch the session stats line at the bottom of the Web UI and the per-turn end line (time / first token / tok/s); estimate how much of this task hit cache
- > Check yourself: the 14.7 measurement methods and the 6.3 stats-line format
3. **Hands-on**: run the same batch of tasks two ways — "a new session per item" vs "all in one session" — and compare wall-clock times; the new-session version will almost certainly be noticeably slower on the first turn (cold start)
- > Check yourself: 14.1's measured 110s cold start vs 1s hot cache

4. **Think**: why does "looking only at this turn's total tokens" misjudge cost? Which two numbers did the #735 comment thread suggest splitting the per-turn display into?
- > Check yourself: the "Community suggestion" paragraph in 14.7
5. **Hands-on**: using #735's suggested metric, manually split one session's total tokens into "hit/miss" parts, price them with 14.4's unit prices, and verify the "order of magnitude" claim
- > Check yourself: the 14.4 comparison table and the 14.2 measurement table

FAQ

- **Q**: What's the highest possible cache hit rate? 97% measured in this handbook (Section 5.1.1); community long runs reach 99.7% ([#560](https://github.com/deepseek-ai/deepseek-harness/discussions/560)). But **brand-new tasks / cold starts drop noticeably** (Section 12.5) — the hit rate is a function of *how you work*, not a constant of the model.
- **Q**: Does the `low` tier noticeably hurt quality? Almost no difference for simple deterministic tasks (file ops, batch processing); complex reasoning, long-chain planning, and debugging at `low` can miss critical steps (FAQ Q3 in Chapter 6). Recommended: `high` for everyday, `low` for batch/simple tasks.
- **Q**: When will the official per-turn token display ship? [#735](https://github.com/deepseek-ai/deepseek-harness/discussions/735) is a feature request filed 2026-08-14 (2 comments) — **not implemented yet**. Until then: the session stats line's hit % + the community `dsh-plugin-cost-tracker` (Chapter 9).
- **Q**: Do third-party gateways / self-hosted providers get the cache discount too? **The cache discount is a DeepSeek API-side capability** (Chapter 5 pricing table); third-party gateways' cache support and billing must be checked on their own pricing pages (inferred, pending verification).
- **Q**: Why does a task get noticeably faster after a cache hit? The hit prefix doesn't need recomputing, so first-token time drops by an order of magnitude (cold start ~110s vs hot cache ~1s, measured in Chapter 1). Tell-tale sign: first token suddenly drops from 10s+ to under 1s — almost certainly a hit (FAQ Q4 in Chapter 6). Use fresh prompts for performance comparisons to avoid cache interference (Chapter 6 pitfall #6).

Info as of 2026-08-14 (dsh 0.1.0-rc.6). Pricing and cache policy are subject to the official docs; items marked "(inferred, pending verification)" are extrapolations or estimates — corrections via PR after real-world testing are welcome.

[English](./15-ecosystem-report.en.md) | [中文](./15-ecosystem-report.md) · [[← Back](#)]
(../README.md)

Chapter 15: The dsh Ecosystem Panorama — What 1,800+ Plugins and 780 Discussion Posts Tell Us

> **Goal of this chapter:** Cross-validate the **community plugin inventory (1,804 plugins)** against **this handbook's 780-post discussion-board observations plus 195-post response experience**, and draw the panorama of the dsh ecosystem — where it's booming, where the gaps are, and where the opportunities lie.

> **Data sources:** the Blue-Whale-Harness plugin master list (community-maintained; 1,804 repos crawled 2026-08, of which 1,663 are flagged "true DSH"); discussion-board observations as of 2026-08.

TL;DR (30-second version)

- 1. **The ecosystem is exploding**: 1,804 plugin repos grown from zero within two weeks; design/UI generation (348 / 94K★) and vision (132 / 90K★) are the hottest demand
 - 2. **"Whatever the official side didn't build, the community did" is the theme**: desktop shells (140+), TUIs, memory, vision bridges — all filled in by the community
 - 3. **Clear gaps**: finance (8), database (14), audio (19), evaluation (37) are nearly empty — yet the demand side exists
 - 4. **Windows and remote are the sore spot**: the Chinese-path family (#107 and 15+ posts), subprocess (#717), LAN RPC (#755)
 - 5. **The differentiation bet**: full session logs + everything-is-a-plugin + sandbox/permission system → enterprise / evaluation / compliance scenarios
-

15.1 Ecosystem Data Snapshot (1,804 repos)

Category Distribution

Category	Plugins	Read	
tools	569	Tools are the hottest category the core value carrier of an agent	
utility	345	General-purpose utilities	

session	229	Session management logs / recovery / statistics	
orchestration	197	Orchestration / workflows	
ui	176	UI enhancements sidebar / panels / components	
llm	159	Model adaptation	
acp	66	ACP protocol	
skills	51	Skill packs	
sandbox	9	Only 9 sandbox plugins — security tooling is scarce see 15.4	

Theme Heat (plugins / total stars / read)

Theme	Plugins	Total ★	Read	
Design / UI generation	348	94K	Hottest demand agents producing visuals is a hard need	
IDE integration	221	89K	High demand	
Vision / image reading	132	90K	Text-only models reading images is a hard need corroborated by the vision-bridge family in #1153 / #1269 / #1378	
Desktop / clients	140	90K	No official desktop shell; the community filled the gap Windows / macOS / Android	
Content creation	40	88K	High value per project few projects but high stars — most	

			valuable for non-programmer scenarios	
Workflow / orchestration	94	17K	Medium	
Memory / persona	77	7.7K	Gaining traction #1822 memory core / #1448 / #1478 etc., many routes	
Code quality / review	65	1.7K	Medium	
Finance / trading	8	28	Gap	
Database	14	137	Gap	
Multimodal / audio	19	26	Gap	
Evaluation / testing	37	117	Gap contrasts with the official 100%-coverage gate	

15.2 Qualitative Insights (780-post discussion-board observation, cross-validated with data)

Insight 1: Windows is the #1 pain point (60+ posts)

Data side: Windows compatibility is the highest-frequency topic; discussion side:

****Chinese-path truncation (#107 / #151 / #396 / #800 and 15+ posts with the same root cause)****, koffi native crashes (#197 / #293), reserved port ranges (#589), subprocesses without graceful termination (#717), sandbox temp-dir cleanup crashes (#758).

****Cross-validation****: the data side's "poor Windows compatibility" and the discussion side's 60+ posts fully corroborate each other — ****this is the most certain ecosystem shortcoming****.

Insight 2: "Whatever the official side didn't build, the community did"

Data side: 140+ desktop shells, multiple TUIs, 77 memory, 132 vision; discussion side: within two weeks the community filled in CLI / TUI / desktop / memory / vision (#391 / #405 / #386 etc.).

****Cross-validation****: the official side (focused on the core during 0.1.0-rc) and the community (filling niche positions) form a healthy complement — ****the strongest signal of ecosystem vitality****.

Insight 3: Security auditing is unusually active

Data side: ****only 9 plugins in the sandbox category**** (security tooling is scarce); discussion side: dense security-audit threads (#243 / #250 / #278 / #381 / #454 / #774 / #817 systematic audits; vm escape / approval loops / clickjacking — real findings) plus a security toolchain (vetting / egress-guard / read-confine / check-dsh-profile).

****Cross-validation****: security demand is high (active auditing) but security tools are few (9 sandbox plugins) — ****a supply gap****.

Insight 4: A family of serialization / boundary bugs is breaking out at once

Discussion side: unknown tool "" (#725 / #1405), reasoning serialization elision (#739 / #1850), truncated tool-call contamination (#1519), run_code async callback drops (#1476), corrupted sessions dragging down boot (#1473 / #1497) — ****all boundary issues in the llm-deepseek serialization layer / session state machine****.

****Read****: this cluster of bugs concentrated in the rc period = the official side's next iteration battleground, and a goldmine for community patches / fix branches (#1519 / #725 both ship with test-backed fixes).

Insight 5: Cost transparency is a hidden must-have

Data side: content creation 40 / 88K★; discussion side: dense token-consumption / cache-hit discussions (#735 / #1052 / #1234 / #1374 / #1520), cost tools springing up like mushrooms (dsh-usage / usage-panel / Show Me How Much / balance-meter).

****Cross-validation****: a 97% cache hit rate (measured) is a cost advantage, but it's "invisible" — cost tooling moving from "statistics" to "engineering" (Chapter 14) is a certain direction.

15.3 Capability Gaps the Official Side Could Prioritize (double-verified by data + discussion)

Capability gap	Data basis	Discussion basis	Feasibility	
Multimodal / vision channel OCR + vision-service abstraction	132 vision / 90K★	#1378 / #1327 / #2024 raised repeatedly	High a core implementation is already emerging	
Memory / persona seam officially replaceable	77 memory / 7.7K★ and gaining traction	#1822 memory-core proposal	High session already keeps full records	
Desktop / TUI shell protocol standardize and absorb	140+ fragmented desktop shells	Multiple coexisting clients	High	
Evaluation / verification loop session → report	Only 37 eval / 117★	Official 100%-coverage gate culture	Medium-high	
First-class Windows support subprocess / sandbox / paths / remote	Windows is the highest-frequency topic	#107 family / #717 / #755	Medium-high large engineering effort	
Plugin registry + compatibility contract	1,804 repos need distribution	#1825 plugin-market meta-plugin	Medium	

15.4 Advice for Ecosystem Participants

If you're a plugin developer

1. ****Prioritize the gaps****: finance / database / audio / evaluation are blue ocean (few competitors, confirmed demand)
2. ****Security tooling is scarce****: only 9 plugins in the sandbox category — security / permission / audit plugins are under-supplied

3. **Vision / memory are hotspots**: but competition is fierce — differentiation (vertical scenarios / enterprise) beats following the crowd

If you're evaluating options

1. **The differentiation bet**: full session logs + everything-is-a-plugin + sandbox/permissions → **enterprise / evaluation / compliance scenarios** (the advantage zone double-verified by data and discussion)
2. **Risk points**: Windows / remote scenarios are currently the sore spot; teams primarily on Windows need to evaluate carefully
3. **Cost advantage**: 97% cache hit × discount, and the cost-transparency toolchain is already complete (Chapter 14)

If you're watching from the sidelines

- **The early-ecosystem dividend is still here**: 1,804 plugins, but vertical domains (finance 8 / database 14 / audio 19) are nearly empty — the window is clear
- **First movers, be patient**: the official rc iteration is fast; capability gaps will keep appearing

15.5 Methodology Notes for This Report

- The data side (1,804 repos) comes from the community master list, with the scope annotated (README mentions dsh, includes some non-true plugins) — **a directional reference, not precise statistics**
 - The qualitative side (780 posts + 195-post responses) comes from this handbook's ongoing observation — every post number is traceable
 - The two sides are compared, and conclusions take "what both data and discussion point to", reducing one-sided bias
- > The ecosystem is changing; this report will keep updating along with the discussion board (see [feedback pipeline](./research/feedback-pipeline.md)).*